

# Model\_Scratch 01 - Data augmentation

Name: Alberto Pingo

Email: 2202145 @my.ipleiria.pt

Validation dataset: train5

## Directories

This section sets up the directory paths used for training, validation, and test datasets based on the repository structure.

```
In [ ]: import os

current_dir = os.getcwd()

# TWO FOLDERS UP
data_dir = os.path.abspath(os.path.join(current_dir, os.pardir, os.pardir))
test_dir = os.path.join(data_dir, 'test')
train_dir = os.path.join(data_dir, 'train')

train_dirs = []
for i in range(1, 5):
    train_dirs.append(os.path.join(train_dir, 'train' + str(i)))

validation_dir = os.path.join(data_dir, 'train', 'train5')

print(current_dir)
print(data_dir)
print(test_dir)
print(train_dir)
print(validation_dir)
```

```
/home/pws/code/IA-image-classification/notebooks/models-S
/home/pws/code/IA-image-classification/data
/home/pws/code/IA-image-classification/data/test
/home/pws/code/IA-image-classification/data/train
/home/pws/code/IA-image-classification/data/train/train5
```

## Preprocessing

Load the datasets and perform initial preprocessing. Images are resized to 32x32 pixels and batched.

```
In [ ]: from keras.utils import image_dataset_from_directory
import tensorflow as tf

# Load training datasets from train1 to train4
train_datasets = []
IMG_SIZE = 32
BATCH_SIZE = 64
```

```
for i in range(1, 5):
    dataset = image_dataset_from_directory(train_dirs[i-1], image_size=(I
    train_datasets.append(dataset)

train_dataset = train_datasets[0]
for dataset in train_datasets[1:]:
    train_dataset = train_dataset.concatenate(dataset)

validation_dataset = image_dataset_from_directory(validation_dir, image_s

test_dataset = image_dataset_from_directory(test_dir, image_size=(IMG_SIZ

class_names = validation_dataset.class_names
class_names = [class_name.split('_')[-1] for class_name in class_names]

print(class_names)
```

2024-06-22 20:44:22.995748: E external/local\_xla/xla/stream\_executor/cuda/cuda\_dnn.cc:9261] Unable to register cuDNN factory: Attempting to register factory for plugin cuDNN when one has already been registered

2024-06-22 20:44:22.995780: E external/local\_xla/xla/stream\_executor/cuda/cuda\_fft.cc:607] Unable to register cuFFT factory: Attempting to register factory for plugin cuFFT when one has already been registered

2024-06-22 20:44:22.996801: E external/local\_xla/xla/stream\_executor/cuda/cuda\_blas.cc:1515] Unable to register cuBLAS factory: Attempting to register factory for plugin cuBLAS when one has already been registered

2024-06-22 20:44:23.002659: I tensorflow/core/platform/cpu\_feature\_guard.cc:182] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.

To enable the following instructions: AVX2 FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.

2024-06-22 20:44:24.090481: W tensorflow/compiler/tf2tensorrt/utils/py\_utils.cc:38] TF-TRT Warning: Could not find TensorRT

Found 10000 files belonging to 10 classes.

```
2024-06-22 20:44:25.464478: I external/local_xla/xla/stream_executor/cuda/cuda_executor.cc:901] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero. See more at https://github.com/torvalds/linux/blob/v6.0/Documentation/ABI/testing/sysfs-bus-pci#L344-L355
2024-06-22 20:44:25.496996: I external/local_xla/xla/stream_executor/cuda/cuda_executor.cc:901] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero. See more at https://github.com/torvalds/linux/blob/v6.0/Documentation/ABI/testing/sysfs-bus-pci#L344-L355
2024-06-22 20:44:25.497164: I external/local_xla/xla/stream_executor/cuda/cuda_executor.cc:901] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero. See more at https://github.com/torvalds/linux/blob/v6.0/Documentation/ABI/testing/sysfs-bus-pci#L344-L355
2024-06-22 20:44:25.497675: I external/local_xla/xla/stream_executor/cuda/cuda_executor.cc:901] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero. See more at https://github.com/torvalds/linux/blob/v6.0/Documentation/ABI/testing/sysfs-bus-pci#L344-L355
2024-06-22 20:44:25.497802: I external/local_xla/xla/stream_executor/cuda/cuda_executor.cc:901] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero. See more at https://github.com/torvalds/linux/blob/v6.0/Documentation/ABI/testing/sysfs-bus-pci#L344-L355
2024-06-22 20:44:25.497931: I external/local_xla/xla/stream_executor/cuda/cuda_executor.cc:901] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero. See more at https://github.com/torvalds/linux/blob/v6.0/Documentation/ABI/testing/sysfs-bus-pci#L344-L355
2024-06-22 20:44:25.565144: I external/local_xla/xla/stream_executor/cuda/cuda_executor.cc:901] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero. See more at https://github.com/torvalds/linux/blob/v6.0/Documentation/ABI/testing/sysfs-bus-pci#L344-L355
2024-06-22 20:44:25.565298: I external/local_xla/xla/stream_executor/cuda/cuda_executor.cc:901] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero. See more at https://github.com/torvalds/linux/blob/v6.0/Documentation/ABI/testing/sysfs-bus-pci#L344-L355
2024-06-22 20:44:25.565427: I external/local_xla/xla/stream_executor/cuda/cuda_executor.cc:901] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero. See more at https://github.com/torvalds/linux/blob/v6.0/Documentation/ABI/testing/sysfs-bus-pci#L344-L355
2024-06-22 20:44:25.565525: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1929] Created device /job:localhost/replica:0/task:0/device:GPU:0 with 299 MB memory: -> device: 0, name: NVIDIA GeForce GTX 1060 6GB, pci bus id: 0000:01:00.0, compute capability: 6.1
Found 10000 files belonging to 10 classes.
Found 10000 files belonging to 10 classes.
Found 10000 files belonging to 10 classes.
Found 10000 files belonging to 10 classes.
Found 10000 files belonging to 10 classes.
['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
```

Configure the dataset for performance

```
In [ ]: AUTOTUNE = tf.data.AUTOTUNE

train_dataset = train_dataset.cache().shuffle(1000).prefetch(buffer_size=
validation_dataset = validation_dataset.cache().prefetch(buffer_size=AUTO
test_dataset = test_dataset.cache().prefetch(buffer_size=AUTOTUNE)
```

## Data Augmentation

Random change of flipping the image horizontally.

Random chance of moving the image horizontally and vertically [-10%, 10%].

Tried with a more complex approach to data augmentation, but the results were worse because of the small size of the images.

```
In [ ]: from keras import layers

data_augmentation = tf.keras.Sequential([
    layers.RandomFlip("horizontal"),
    layers.RandomTranslation(0.1, 0.1, fill_mode='nearest'),
])



In [ ]: import matplotlib.pyplot as plt
import numpy as np



#Plot some Augmented images
for images, labels in train_dataset.take(1):
    plt.figure(figsize=(10, 10))
    first_image = images[0]
    for i in range(4):
        ax = plt.subplot(2, 2, i + 1)
        augmented_image = data_augmentation(tf.expand_dims(first_image, 0)
        plt.imshow(augmented_image[0] / 255)
        plt.axis('off')
```





```
In [ ]: import matplotlib.pyplot as plt



for data, _ in train_dataset.take(1):
    for i in range(9):
        plt.imshow(data[i].numpy().astype('uint8'))
        plt.show()
    break
```














## MODEL ARCHITECTURE

### Build a Convolutional Neural Network (CNN) model.

#### Architecture:

Input -> Conv2D - BN-> Conv2D - BN-> MaxPooling2D -> Conv2D - BN-> Conv2D - BN-> MaxPooling2D -> Flatten -> Dense -> Dropout -> Dense -> Dropout -> Output

#### 1. Input Layer

- The input layer expects images of size 32x32 pixels with 3 color channels (RGB).
- Data augmentation is applied to the input images.
- The Rescaling layer, rescales the pixel values from the range [0, 255] to [0, 1].

#### 2. Convolutional Layers

- The model consists of 4 convolutional layers with 32, 64, 128, and 128 filters respectively.

#### 3. Max Pooling Layers

- Max pooling layers are used after each group of 2 convolutional layer to reduce the spatial dimensions of the feature maps.
- A pooling size of 2x2 is used.

#### 4. Fully connected layer

- A dense layer with 512 units and ReLU activation function.

## 5. Output Layer

- The output layer consists of 10 units (one for each class) with a softmax activation function.
- The softmax function outputs the probability distribution over the classes.

## Overfitting measures

- Dropout layers are used after each Convolutional and Dense layer to prevent overfitting.
- Kernel Regularization is used to prevent overfitting.

## Batch Normalization

- Batch normalization is used after each Convolutional layer to normalize the activations of the previous layer at each batch.
- This helps to stabilize and speed up the training process.

## Weight Initialization

- For the ReLU activation function, the `he_normal` initializer used to initialize the weights is considered a good weight initialization for ReLU activation functions.
- For the output layer using the `Softmax` activation function, the `glorot_uniform` initializer is used for the same reason.

```
In [ ]: from tensorflow import keras
        from keras import layers, regularizers

        inputs = keras.Input(shape=(IMG_SIZE, IMG_SIZE, 3))

        x = data_augmentation(inputs)

        x = layers.Rescaling(1./255)(x)

        ## First Convolutional Block
        x = layers.Conv2D(filters=32, kernel_size=3, kernel_initializer='he_normal')(x)
        x = layers.BatchNormalization()(x) # Standardize the inputs to the next l
        x = layers.Conv2D(filters=64, kernel_size=3, kernel_initializer='he_normal')(x)
        x = layers.BatchNormalization()(x)

        # First Block - Max Pooling and Dropout
        x = layers.MaxPooling2D(pool_size=2)(x)
        x = layers.Dropout(0.3)(x) # Drops 30% of the neurons randomly during tra

        # Second Convolutional Block
        x = layers.Conv2D(filters=128, kernel_size=3, kernel_initializer='he_normal')(x)
        x = layers.BatchNormalization()(x)
        x = layers.Conv2D(filters=128, kernel_size=3, kernel_initializer='he_normal')(x)
        x = layers.BatchNormalization()(x)

        # Second Block - Max Pooling and Dropout
        x = layers.MaxPooling2D(pool_size=2)(x)
        x = layers.Dropout(0.3)(x)

        x = layers.Flatten()(x)
```

```
x = layers.Dense(512, kernel_initializer='he_normal', activation="relu")(
x = layers.Dropout(0.5)(x)

outputs = layers.Dense(10, kernel_initializer='glorot_uniform', activati

model = keras.Model(inputs=inputs, outputs=outputs)
model.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 32, 32, 3)]	0
sequential (Sequential)	(None, 32, 32, 3)	0
rescaling (Rescaling)	(None, 32, 32, 3)	0
conv2d (Conv2D)	(None, 30, 30, 32)	896
batch_normalization (Batch Normalization)	(None, 30, 30, 32)	128
conv2d_1 (Conv2D)	(None, 28, 28, 64)	18496
batch_normalization_1 (Batch Normalization)	(None, 28, 28, 64)	256
max_pooling2d (MaxPooling2D)	(None, 14, 14, 64)	0
dropout (Dropout)	(None, 14, 14, 64)	0
conv2d_2 (Conv2D)	(None, 12, 12, 128)	73856
batch_normalization_2 (Batch Normalization)	(None, 12, 12, 128)	512
conv2d_3 (Conv2D)	(None, 10, 10, 128)	147584

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 32, 32, 3)]	0
sequential (Sequential)	(None, 32, 32, 3)	0
rescaling (Rescaling)	(None, 32, 32, 3)	0
conv2d (Conv2D)	(None, 30, 30, 32)	896
batch_normalization (Batch Normalization)	(None, 30, 30, 32)	128
conv2d_1 (Conv2D)	(None, 28, 28, 64)	18496
batch_normalization_1 (Batch Normalization)	(None, 28, 28, 64)	256
max_pooling2d (MaxPooling2D)	(None, 14, 14, 64)	0
dropout (Dropout)	(None, 14, 14, 64)	0
conv2d_2 (Conv2D)	(None, 12, 12, 128)	73856
batch_normalization_2 (Batch Normalization)	(None, 12, 12, 128)	512

conv2d_3 (Conv2D)	(None, 10, 10, 128)	147584
batch_normalization_3 (Batch Normalization)	(None, 10, 10, 128)	512
max_pooling2d_1 (MaxPooling2D)	(None, 5, 5, 128)	0
dropout_1 (Dropout)	(None, 5, 5, 128)	0
flatten (Flatten)	(None, 3200)	0
dense (Dense)	(None, 512)	1638912
dropout_2 (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 10)	5130

---

=====  
 Total params: 1886282 (7.20 MB)  
 Trainable params: 1885578 (7.19 MB)  
 Non-trainable params: 704 (2.75 KB)

---

## Compile Model

### Loss function:

We use the *Categorical Crossentropy* loss function because it is a multi-class classification problem.

### Optimizer: Adam

We use the *Adam* optimizer because it is one of the best and most popular optimizers.

```
In [ ]: model.compile(
        loss='categorical_crossentropy',
        optimizer='adam',
        metrics=['acc'])
```

## Train Model

Train the model with Early stopping, Model checkpoint, and Learning rate reduction callbacks.

```
In [ ]: from keras.callbacks import EarlyStopping, ModelCheckpoint, ReduceLRonPlateau

learning_rate_reduction = ReduceLRonPlateau(
    monitor='val_acc',
    patience=3,
    verbose=1,
    factor=0.5,
    min_lr=1e-5)

early_stop = EarlyStopping(monitor='val_acc',
                           patience=7,
```

```

                                restore_best_weights=True)
model_checkpoint = ModelCheckpoint('models/S01/checkpoints/S01-DA-cp.h5',

history = model.fit(
    train_dataset,
    epochs=100,
    validation_data=validation_dataset,
    callbacks=[early_stop, model_checkpoint, learning_rate_reduction])

```

Epoch 1/100

```

2024-06-22 20:44:33.120120: E tensorflow/core/grappler/optimizers/meta_opt
imizer.cc:961] layout failed: INVALID_ARGUMENT: Size of values 0 does not
match size of permutation 4 @ fanin shape inmodel/dropout/dropout/SelectV
2-2-TransposeNHWCToNCHW-LayoutOptimizer
2024-06-22 20:44:33.318379: I external/local_xla/xla/stream_executor/cuda/
cuda_dnn.cc:454] Loaded cuDNN version 8904
2024-06-22 20:44:34.433995: I external/local_xla/xla/service/service.cc:16
8] XLA service 0x7947f4ac3560 initialized for platform CUDA (this does not
guarantee that XLA will be used). Devices:
2024-06-22 20:44:34.434025: I external/local_xla/xla/service/service.cc:17
6] StreamExecutor device (0): NVIDIA GeForce GTX 1060 6GB, Compute Capab
ility 6.1
2024-06-22 20:44:34.439490: I tensorflow/compiler/mlir/tensorflow/utils/du
mp_mlir_util.cc:269] disabling MLIR crash reproducer, set env var `MLIR_CR
ASH_REPRODUCER_DIRECTORY` to enable.
WARNING: All log messages before absl::InitializeLog() is called are writt
en to STDERR
I0000 00:00:1719085474.519916 120032 device_compiler.h:186] Compiled clus
ter using XLA! This line is logged at most once for the lifetime of the p
rocess.
628/628 [=====] - 18s 22ms/step - loss: 1.8973 -
acc: 0.3641 - val_loss: 1.9815 - val_acc: 0.4068 - lr: 0.0010

```

Epoch 2/100

```

  4/628 [.....] - ETA: 13s - loss: 1.5817 - acc:
0.4805

```

```

/home/pws/miniconda3/envs/tensorflow/lib/python3.11/site-packages/keras/sr
c/engine/training.py:3103: UserWarning: You are saving your model as an HD
F5 file via `model.save()`. This file format is considered legacy. We reco
mmend using instead the native Keras format, e.g. `model.save('my_model.ke
ras')`.

```

```

    saving_api.save_model(

```

```
628/628 [=====] - 13s 20ms/step - loss: 1.4920 -  
acc: 0.4936 - val_loss: 1.3209 - val_acc: 0.5680 - lr: 0.0010  
Epoch 3/100  
628/628 [=====] - 13s 21ms/step - loss: 1.3165 -  
acc: 0.5602 - val_loss: 1.1575 - val_acc: 0.6255 - lr: 0.0010  
Epoch 4/100  
628/628 [=====] - 13s 21ms/step - loss: 1.1937 -  
acc: 0.6095 - val_loss: 1.0305 - val_acc: 0.6679 - lr: 0.0010  
Epoch 5/100  
628/628 [=====] - 13s 21ms/step - loss: 1.1123 -  
acc: 0.6425 - val_loss: 1.0317 - val_acc: 0.6793 - lr: 0.0010  
Epoch 6/100  
628/628 [=====] - 13s 20ms/step - loss: 1.0571 -  
acc: 0.6645 - val_loss: 0.9111 - val_acc: 0.7095 - lr: 0.0010  
Epoch 7/100  
628/628 [=====] - 13s 20ms/step - loss: 1.0131 -  
acc: 0.6822 - val_loss: 0.9219 - val_acc: 0.7183 - lr: 0.0010  
Epoch 8/100  
628/628 [=====] - 13s 20ms/step - loss: 0.9734 -  
acc: 0.7006 - val_loss: 0.8070 - val_acc: 0.7564 - lr: 0.0010  
Epoch 9/100  
628/628 [=====] - 13s 21ms/step - loss: 0.9382 -  
acc: 0.7161 - val_loss: 0.7755 - val_acc: 0.7641 - lr: 0.0010  
Epoch 10/100  
628/628 [=====] - 13s 21ms/step - loss: 0.9213 -  
acc: 0.7214 - val_loss: 0.8992 - val_acc: 0.7411 - lr: 0.0010  
Epoch 11/100  
628/628 [=====] - 13s 21ms/step - loss: 0.8991 -  
acc: 0.7330 - val_loss: 0.7490 - val_acc: 0.7832 - lr: 0.0010  
Epoch 12/100  
628/628 [=====] - 13s 20ms/step - loss: 0.8784 -  
acc: 0.7376 - val_loss: 0.7693 - val_acc: 0.7748 - lr: 0.0010  
Epoch 13/100  
628/628 [=====] - 13s 21ms/step - loss: 0.8624 -  
acc: 0.7461 - val_loss: 0.7865 - val_acc: 0.7711 - lr: 0.0010  
Epoch 14/100  
626/628 [=====>.] - ETA: 0s - loss: 0.8417 - acc:  
0.7558  
Epoch 14: ReduceLROnPlateau reducing learning rate to 0.000500000023748725  
7.  
628/628 [=====] - 13s 21ms/step - loss: 0.8416 -  
acc: 0.7560 - val_loss: 0.8996 - val_acc: 0.7381 - lr: 0.0010  
Epoch 15/100  
628/628 [=====] - 13s 21ms/step - loss: 0.7699 -  
acc: 0.7773 - val_loss: 0.7265 - val_acc: 0.7888 - lr: 5.0000e-04  
Epoch 16/100  
628/628 [=====] - 13s 21ms/step - loss: 0.7379 -  
acc: 0.7888 - val_loss: 0.6861 - val_acc: 0.8112 - lr: 5.0000e-04  
Epoch 17/100  
628/628 [=====] - 13s 21ms/step - loss: 0.7204 -  
acc: 0.7926 - val_loss: 0.6775 - val_acc: 0.8102 - lr: 5.0000e-04  
Epoch 18/100  
628/628 [=====] - 13s 20ms/step - loss: 0.7026 -  
acc: 0.7981 - val_loss: 0.6817 - val_acc: 0.8063 - lr: 5.0000e-04  
Epoch 19/100  
628/628 [=====] - 13s 21ms/step - loss: 0.6887 -  
acc: 0.8020 - val_loss: 0.6258 - val_acc: 0.8214 - lr: 5.0000e-04  
Epoch 20/100  
628/628 [=====] - 13s 21ms/step - loss: 0.6844 -  
acc: 0.8029 - val_loss: 0.6144 - val_acc: 0.8290 - lr: 5.0000e-04
```



```
Epoch 21/100
628/628 [=====] - 13s 21ms/step - loss: 0.6706 -
acc: 0.8067 - val_loss: 0.6152 - val_acc: 0.8281 - lr: 5.0000e-04
Epoch 22/100
628/628 [=====] - 13s 21ms/step - loss: 0.6621 -
acc: 0.8098 - val_loss: 0.5726 - val_acc: 0.8431 - lr: 5.0000e-04
Epoch 23/100
628/628 [=====] - 13s 20ms/step - loss: 0.6544 -
acc: 0.8130 - val_loss: 0.6520 - val_acc: 0.8167 - lr: 5.0000e-04
Epoch 24/100
628/628 [=====] - 13s 20ms/step - loss: 0.6410 -
acc: 0.8162 - val_loss: 0.5887 - val_acc: 0.8355 - lr: 5.0000e-04
Epoch 25/100
627/628 [=====>.] - ETA: 0s - loss: 0.6370 - acc:
0.8180
Epoch 25: ReduceLRonPlateau reducing learning rate to 0.000250000011874362
8.
628/628 [=====] - 13s 21ms/step - loss: 0.6369 -
acc: 0.8180 - val_loss: 0.5882 - val_acc: 0.8371 - lr: 5.0000e-04
Epoch 26/100
628/628 [=====] - 13s 21ms/step - loss: 0.5971 -
acc: 0.8313 - val_loss: 0.5663 - val_acc: 0.8429 - lr: 2.5000e-04
Epoch 27/100
628/628 [=====] - 13s 20ms/step - loss: 0.5829 -
acc: 0.8339 - val_loss: 0.5778 - val_acc: 0.8435 - lr: 2.5000e-04
Epoch 28/100
628/628 [=====] - 13s 20ms/step - loss: 0.5661 -
acc: 0.8408 - val_loss: 0.5358 - val_acc: 0.8544 - lr: 2.5000e-04
Epoch 29/100
628/628 [=====] - 13s 20ms/step - loss: 0.5622 -
acc: 0.8427 - val_loss: 0.5450 - val_acc: 0.8526 - lr: 2.5000e-04
Epoch 30/100
628/628 [=====] - 13s 20ms/step - loss: 0.5555 -
acc: 0.8437 - val_loss: 0.5246 - val_acc: 0.8578 - lr: 2.5000e-04
Epoch 31/100
628/628 [=====] - 13s 20ms/step - loss: 0.5489 -
acc: 0.8464 - val_loss: 0.5552 - val_acc: 0.8457 - lr: 2.5000e-04
Epoch 32/100
628/628 [=====] - 13s 20ms/step - loss: 0.5427 -
acc: 0.8483 - val_loss: 0.5598 - val_acc: 0.8423 - lr: 2.5000e-04
Epoch 33/100
626/628 [=====>.] - ETA: 0s - loss: 0.5355 - acc:
0.8499
Epoch 33: ReduceLRonPlateau reducing learning rate to 0.000125000005937181
4.
628/628 [=====] - 13s 21ms/step - loss: 0.5357 -
acc: 0.8499 - val_loss: 0.5237 - val_acc: 0.8559 - lr: 2.5000e-04
Epoch 34/100
628/628 [=====] - 13s 21ms/step - loss: 0.5160 -
acc: 0.8563 - val_loss: 0.5276 - val_acc: 0.8570 - lr: 1.2500e-04
Epoch 35/100
628/628 [=====] - 13s 20ms/step - loss: 0.5112 -
acc: 0.8562 - val_loss: 0.5252 - val_acc: 0.8607 - lr: 1.2500e-04
Epoch 36/100
628/628 [=====] - 12s 18ms/step - loss: 0.5016 -
acc: 0.8611 - val_loss: 0.5083 - val_acc: 0.8611 - lr: 1.2500e-04
Epoch 37/100
628/628 [=====] - 11s 17ms/step - loss: 0.4921 -
acc: 0.8616 - val_loss: 0.5147 - val_acc: 0.8619 - lr: 1.2500e-04
Epoch 38/100
```

```
628/628 [=====] - 11s 17ms/step - loss: 0.4934 -  
acc: 0.8615 - val_loss: 0.5194 - val_acc: 0.8629 - lr: 1.2500e-04  
Epoch 39/100  
628/628 [=====] - 11s 17ms/step - loss: 0.4887 -  
acc: 0.8630 - val_loss: 0.5109 - val_acc: 0.8621 - lr: 1.2500e-04  
Epoch 40/100  
628/628 [=====] - 11s 17ms/step - loss: 0.4858 -  
acc: 0.8661 - val_loss: 0.5061 - val_acc: 0.8652 - lr: 1.2500e-04  
Epoch 41/100  
628/628 [=====] - 11s 17ms/step - loss: 0.4839 -  
acc: 0.8629 - val_loss: 0.5397 - val_acc: 0.8539 - lr: 1.2500e-04  
Epoch 42/100  
628/628 [=====] - 11s 17ms/step - loss: 0.4867 -  
acc: 0.8645 - val_loss: 0.5156 - val_acc: 0.8573 - lr: 1.2500e-04  
Epoch 43/100  
628/628 [=====] - 11s 17ms/step - loss: 0.4745 -  
acc: 0.8677 - val_loss: 0.5036 - val_acc: 0.8659 - lr: 1.2500e-04  
Epoch 44/100  
628/628 [=====] - 11s 17ms/step - loss: 0.4700 -  
acc: 0.8705 - val_loss: 0.5062 - val_acc: 0.8613 - lr: 1.2500e-04  
Epoch 45/100  
628/628 [=====] - 11s 17ms/step - loss: 0.4676 -  
acc: 0.8687 - val_loss: 0.4858 - val_acc: 0.8707 - lr: 1.2500e-04  
Epoch 46/100  
628/628 [=====] - 11s 17ms/step - loss: 0.4676 -  
acc: 0.8683 - val_loss: 0.4992 - val_acc: 0.8671 - lr: 1.2500e-04  
Epoch 47/100  
628/628 [=====] - 11s 18ms/step - loss: 0.4648 -  
acc: 0.8685 - val_loss: 0.4856 - val_acc: 0.8685 - lr: 1.2500e-04  
Epoch 48/100  
626/628 [=====>.] - ETA: 0s - loss: 0.4582 - acc:  
0.8726  
Epoch 48: ReduceLRonPlateau reducing learning rate to 6.25000029685907e-0  
5.  
628/628 [=====] - 11s 17ms/step - loss: 0.4582 -  
acc: 0.8726 - val_loss: 0.5081 - val_acc: 0.8635 - lr: 1.2500e-04  
Epoch 49/100  
628/628 [=====] - 11s 17ms/step - loss: 0.4534 -  
acc: 0.8744 - val_loss: 0.4932 - val_acc: 0.8663 - lr: 6.2500e-05  
Epoch 50/100  
628/628 [=====] - 11s 17ms/step - loss: 0.4506 -  
acc: 0.8747 - val_loss: 0.4966 - val_acc: 0.8654 - lr: 6.2500e-05  
Epoch 51/100  
628/628 [=====] - 11s 17ms/step - loss: 0.4475 -  
acc: 0.8747 - val_loss: 0.4784 - val_acc: 0.8722 - lr: 6.2500e-05  
Epoch 52/100  
628/628 [=====] - 11s 18ms/step - loss: 0.4476 -  
acc: 0.8760 - val_loss: 0.4831 - val_acc: 0.8691 - lr: 6.2500e-05  
Epoch 53/100  
628/628 [=====] - 11s 17ms/step - loss: 0.4484 -  
acc: 0.8757 - val_loss: 0.4851 - val_acc: 0.8682 - lr: 6.2500e-05  
Epoch 54/100  
628/628 [=====] - ETA: 0s - loss: 0.4421 - acc:  
0.8771  
Epoch 54: ReduceLRonPlateau reducing learning rate to 3.125000148429535e-0  
5.  
628/628 [=====] - 11s 17ms/step - loss: 0.4421 -  
acc: 0.8771 - val_loss: 0.5002 - val_acc: 0.8647 - lr: 6.2500e-05  
Epoch 55/100  
628/628 [=====] - 11s 17ms/step - loss: 0.4303 -
```

```

acc: 0.8819 - val_loss: 0.4880 - val_acc: 0.8683 - lr: 3.1250e-05
Epoch 56/100
628/628 [=====] - 11s 17ms/step - loss: 0.4344 -
acc: 0.8798 - val_loss: 0.4743 - val_acc: 0.8745 - lr: 3.1250e-05
Epoch 57/100
628/628 [=====] - 11s 17ms/step - loss: 0.4310 -
acc: 0.8813 - val_loss: 0.4825 - val_acc: 0.8707 - lr: 3.1250e-05
Epoch 58/100
628/628 [=====] - 11s 18ms/step - loss: 0.4266 -
acc: 0.8827 - val_loss: 0.4868 - val_acc: 0.8692 - lr: 3.1250e-05
Epoch 59/100
626/628 [=====>.] - ETA: 0s - loss: 0.4321 - acc:
0.8809
Epoch 59: ReduceLROnPlateau reducing learning rate to 1.5625000742147677
e-05.
628/628 [=====] - 12s 19ms/step - loss: 0.4320 -
acc: 0.8809 - val_loss: 0.4813 - val_acc: 0.8688 - lr: 3.1250e-05
Epoch 60/100
628/628 [=====] - 11s 18ms/step - loss: 0.4277 -
acc: 0.8825 - val_loss: 0.4780 - val_acc: 0.8702 - lr: 1.5625e-05
Epoch 61/100
628/628 [=====] - 11s 18ms/step - loss: 0.4259 -
acc: 0.8831 - val_loss: 0.4766 - val_acc: 0.8710 - lr: 1.5625e-05
Epoch 62/100
628/628 [=====] - ETA: 0s - loss: 0.4266 - acc:
0.8831
Epoch 62: ReduceLROnPlateau reducing learning rate to 1e-05.
628/628 [=====] - 12s 19ms/step - loss: 0.4266 -
acc: 0.8831 - val_loss: 0.4802 - val_acc: 0.8701 - lr: 1.5625e-05
Epoch 63/100
628/628 [=====] - 11s 18ms/step - loss: 0.4231 -
acc: 0.8838 - val_loss: 0.4805 - val_acc: 0.8711 - lr: 1.0000e-05

```

## Save Model

```
In [ ]: keras.models.save_model(model, 'models/S01/S01-DA-model.h5')
```

```

/tmp/ipykernel_119953/3135399560.py:1: UserWarning: You are saving your mo
del as an HDF5 file via `model.save()`. This file format is considered leg
acy. We recommend using instead the native Keras format, e.g. `model.save
('my_model.keras')`.
  keras.models.save_model(model, 'models/S01/S01-DA-model.h5')

```

## Load Model

```
In [ ]: from tensorflow import keras
model = keras.models.load_model('models/S01/S01-DA-model.h5')
```

## EVALUATION

Evaluate the model on the validation dataset.

```
In [ ]: val_loss, val_acc = model.evaluate(validation_dataset)
        print('val_acc:', val_acc)
```

```
157/157 [=====] - 1s 5ms/step - loss: 0.4743 - ac
c: 0.8745
val_acc: 0.8744999766349792
```

## Training and Validation Curves



Plot the training and validation accuracy and loss curves.

```
In [ ]: import matplotlib.pyplot as plt

# Extract the history from the training process
acc = history.history['acc']
val_acc = history.history['val_acc']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(acc) + 1)

# Plot the training and validation accuracy
plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()

# Plot the training and validation loss
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```



## Confusion Matrix

```
In [ ]: import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix
```

```
y_true = []
y_pred = []

for features, labels in validation_dataset:
    predictions = model.predict(features)
    y_true.extend(np.argmax(labels.numpy(), axis=1))
    y_pred.extend(np.argmax(predictions, axis=1))

y_true = np.array(y_true)
y_pred = np.array(y_pred)

cm = confusion_matrix(y_true, y_pred)


# Plot the confusion matrix
plt.figure(figsize=(10, 8))
sns.heatmap(cm, annot=True, cmap='Blues', fmt='g')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()
```

[illegible]

```
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 3ms/step
2/2 [=====] - 0s 6ms/step
2/2 [=====] - 0s 3ms/step
2/2 [=====] - 0s 3ms/step
2/2 [=====] - 0s 3ms/step
2/2 [=====] - 0s 3ms/step
2/2 [=====] - 0s 3ms/step
2/2 [=====] - 0s 3ms/step
2/2 [=====] - 0s 3ms/step
2/2 [=====] - 0s 4ms/step
2/2 [=====] - 0s 5ms/step
2/2 [=====] - 0s 4ms/step
2/2 [=====] - 0s 3ms/step
2/2 [=====] - 0s 3ms/step
2/2 [=====] - 0s 3ms/step
2/2 [=====] - 0s 3ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 3ms/step
2/2 [=====] - 0s 4ms/step
2/2 [=====] - 0s 4ms/step
2/2 [=====] - 0s 3ms/step
2/2 [=====] - 0s 3ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 3ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 3ms/step
2/2 [=====] - 0s 3ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 4ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 3ms/step
2/2 [=====] - 0s 3ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 3ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 3ms/step
2/2 [=====] - 0s 4ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 2ms/step
```



```
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 3ms/step
2/2 [=====] - 0s 3ms/step
2/2 [=====] - 0s 4ms/step
2/2 [=====] - 0s 3ms/step
2/2 [=====] - 0s 4ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 3ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 4ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 4ms/step
2/2 [=====] - 0s 5ms/step
2/2 [=====] - 0s 10ms/step
2/2 [=====] - 0s 4ms/step
2/2 [=====] - 0s 2ms/step
2/2 [=====] - 0s 3ms/step
2/2 [=====] - 0s 3ms/step
2/2 [=====] - 0s 4ms/step
2/2 [=====] - 0s 2ms/step
1/1 [=====] - 0s 89ms/step
```



```
In [ ]: print("Confusion Matrix:")
        print(cm)
```

Confusion Matrix:

```
[[904 15 23 11 7 1 5 6 27 15]
 [ 2 963 1 1 0 0 1 2 13 31]
 [ 32 2 770 27 32 19 42 22 3 3]
 [ 12 4 27 717 33 101 75 25 6 16]
 [ 10 0 26 19 865 11 41 21 4 0]
 [ 4 1 18 101 35 796 23 40 2 5]
 [ 5 1 10 11 4 6 940 0 3 0]
 [ 6 1 16 13 28 15 6 885 0 7]
 [ 15 14 2 1 2 0 3 0 948 18]
 [ 9 34 2 3 1 0 5 5 6 957]]
```

```
In [ ]: from sklearn.metrics import classification_report

        report = classification_report(y_true, y_pred, target_names=class_names)
        print(report)
```

	precision	recall	f1-score	support
airplane	0.90	0.89	0.90	1014
automobile	0.93	0.95	0.94	1014
bird	0.86	0.81	0.83	952
cat	0.79	0.71	0.75	1016
deer	0.86	0.87	0.86	997
dog	0.84	0.78	0.81	1025
frog	0.82	0.96	0.89	980
horse	0.88	0.91	0.89	977
ship	0.94	0.95	0.94	1003
truck	0.91	0.94	0.92	1022
accuracy			0.87	10000
macro avg	0.87	0.87	0.87	10000
weighted avg	0.87	0.87	0.87	10000

## Predictions

Predict and visualize the results for a sample image.


```
In [ ]: import tensorflow as tf
import matplotlib.pyplot as plt
from keras.preprocessing import image

# Load an image
img = tf.keras.preprocessing.image.load_img(train_dirs[0] + '/006_frog/al
# img = tf.keras.preprocessing.image.load_img(train_dirs[0] + '/000_airpl

img_array = image.img_to_array(img)
img_array = tf.expand_dims(img_array, 0)

plt.imshow(img)
plt.show()

print(img_array.shape)
result = model.predict(img_array)
print("Result: ", result.round())
```



(1, 32, 32, 3)

1/1 [=====] - 0s 229ms/step

Result: [[0. 0. 0. 0. 0. 0. 1. 0. 0. 0.]]