

# Model\_Transfer Learning 01 - No data augmentation - Feature extraction

Name: Alberto Pingo

Email: 2202145 @my.ipleiria.pt

Validation dataset: train5

## Directories

This section sets up the directory paths used for training, validation, and test datasets based on the repository structure.

```
In [ ]: import os

current_dir = os.getcwd()

# TWO FOLDERS UP
data_dir = os.path.abspath(os.path.join(current_dir, os.pardir, os.pardir))
test_dir = os.path.join(data_dir, 'test')
train_dir = os.path.join(data_dir, 'train')

train_dirs = []
for i in range(1, 5):
    train_dirs.append(os.path.join(train_dir, 'train' + str(i)))

validation_dir = os.path.join(data_dir, 'train', 'train5')

print(current_dir)
print(data_dir)
print(test_dir)
print(train_dir)
print(validation_dir)
```

```
/home/pws/code/IA-image-classification/notebooks/models-T
/home/pws/code/IA-image-classification/data
/home/pws/code/IA-image-classification/data/test
/home/pws/code/IA-image-classification/data/train
/home/pws/code/IA-image-classification/data/train/train5
```

## Preprocessing

Load the datasets and perform initial preprocessing. Images are resized to 32x32 pixels and batched.

```
In [ ]: from keras.utils import image_dataset_from_directory
import tensorflow as tf

# Load training datasets from train1 to train4
train_datasets = []
IMG_SIZE = 150
```

```
BATCH_SIZE = 32

for i in range(1, 5):
    dataset = image_dataset_from_directory(train_dirs[i-1], image_size=(I
    train_datasets.append(dataset)

train_dataset = train_datasets[0]
for dataset in train_datasets[1:]:
    train_dataset = train_dataset.concatenate(dataset)

validation_dataset = image_dataset_from_directory(validation_dir, image_s

test_dataset = image_dataset_from_directory(test_dir, image_size=(IMG_SIZ

class_names = validation_dataset.class_names
class_names = [class_name.split('_')[-1] for class_name in class_names]

print(class_names)
```

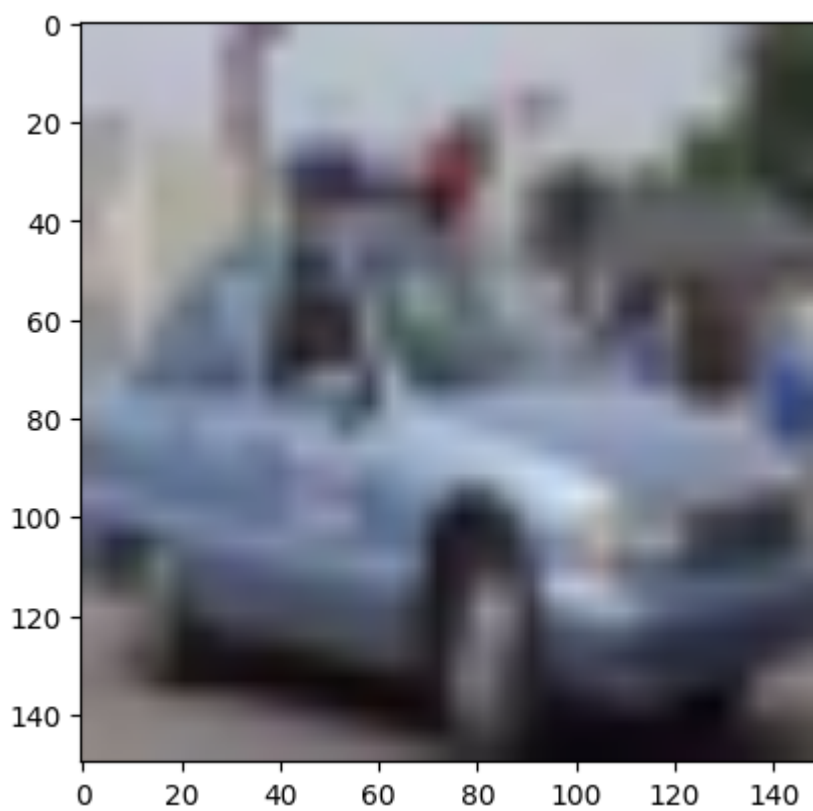
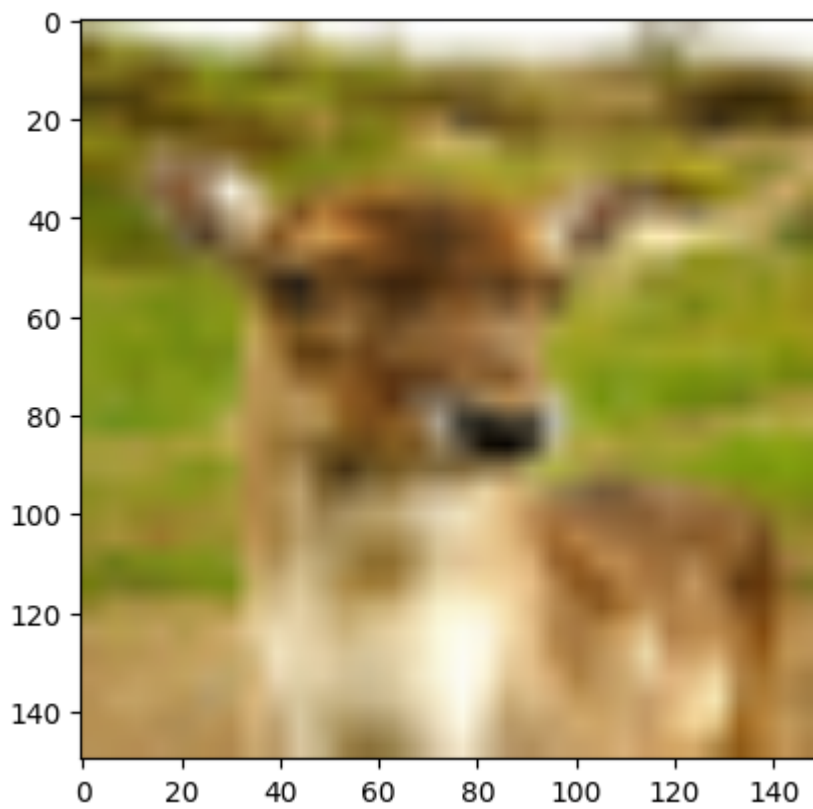
```
2024-06-22 21:29:44.538164: E external/local_xla/xla/stream_executor/cuda/
cuda_dnn.cc:9261] Unable to register cuDNN factory: Attempting to register
factory for plugin cuDNN when one has already been registered
2024-06-22 21:29:44.538220: E external/local_xla/xla/stream_executor/cuda/
cuda_fft.cc:607] Unable to register cuFFT factory: Attempting to register
factory for plugin cuFFT when one has already been registered
2024-06-22 21:29:44.632709: E external/local_xla/xla/stream_executor/cuda/
cuda_blas.cc:1515] Unable to register cuBLAS factory: Attempting to regist
er factory for plugin cuBLAS when one has already been registered
2024-06-22 21:29:44.841271: I tensorflow/core/platform/cpu_feature_guard.c
c:182] This TensorFlow binary is optimized to use available CPU instructio
ns in performance-critical operations.
To enable the following instructions: AVX2 FMA, in other operations, rebui
ld TensorFlow with the appropriate compiler flags.
2024-06-22 21:29:46.151776: W tensorflow/compiler/tf2tensorrt/utils/py_util
s.cc:38] TF-TRT Warning: Could not find TensorRT
Found 10000 files belonging to 10 classes.
```

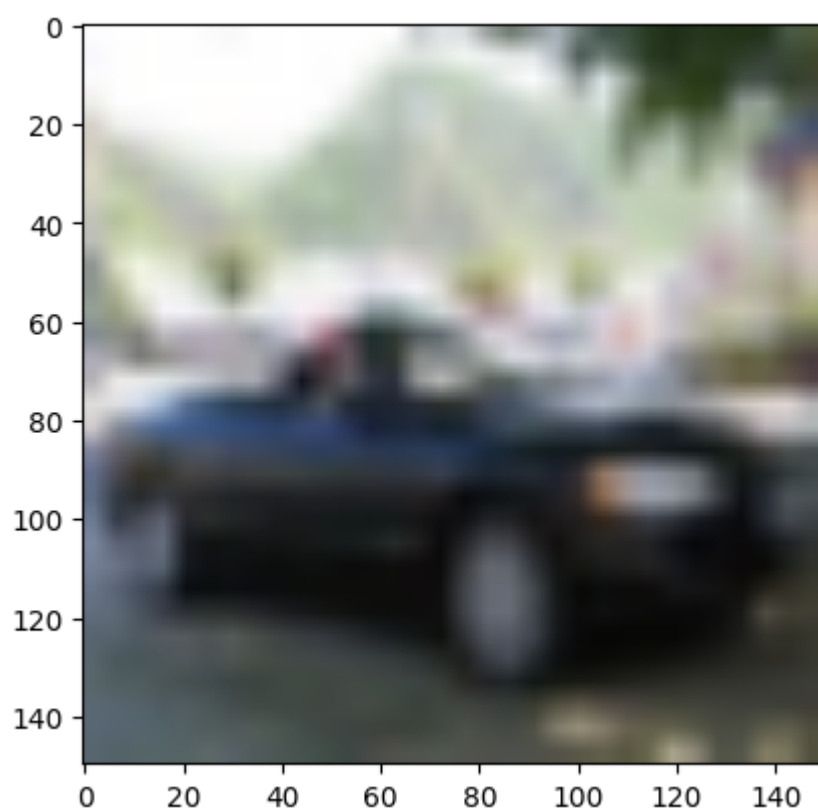
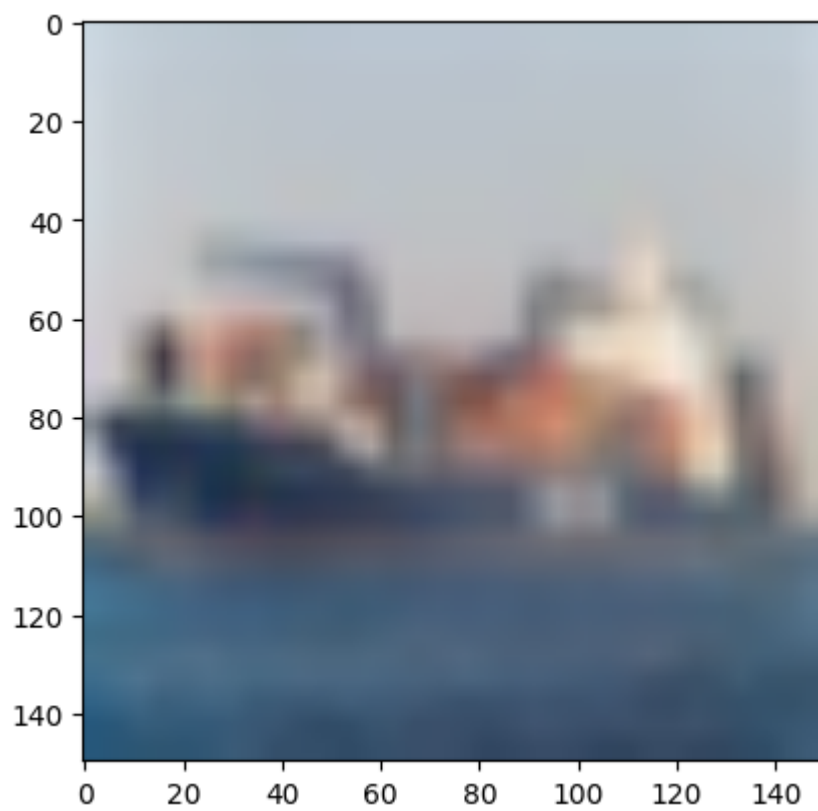
```
2024-06-22 21:29:48.113826: I external/local_xla/xla/stream_executor/cuda/cuda_executor.cc:901] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero. See more at https://github.com/torvalds/linux/blob/v6.0/Documentation/ABI/testing/sysfs-bus-pci#L344-L355
2024-06-22 21:29:48.487821: I external/local_xla/xla/stream_executor/cuda/cuda_executor.cc:901] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero. See more at https://github.com/torvalds/linux/blob/v6.0/Documentation/ABI/testing/sysfs-bus-pci#L344-L355
2024-06-22 21:29:48.488018: I external/local_xla/xla/stream_executor/cuda/cuda_executor.cc:901] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero. See more at https://github.com/torvalds/linux/blob/v6.0/Documentation/ABI/testing/sysfs-bus-pci#L344-L355
2024-06-22 21:29:48.489868: I external/local_xla/xla/stream_executor/cuda/cuda_executor.cc:901] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero. See more at https://github.com/torvalds/linux/blob/v6.0/Documentation/ABI/testing/sysfs-bus-pci#L344-L355
2024-06-22 21:29:48.490057: I external/local_xla/xla/stream_executor/cuda/cuda_executor.cc:901] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero. See more at https://github.com/torvalds/linux/blob/v6.0/Documentation/ABI/testing/sysfs-bus-pci#L344-L355
2024-06-22 21:29:48.490209: I external/local_xla/xla/stream_executor/cuda/cuda_executor.cc:901] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero. See more at https://github.com/torvalds/linux/blob/v6.0/Documentation/ABI/testing/sysfs-bus-pci#L344-L355
2024-06-22 21:29:48.573676: I external/local_xla/xla/stream_executor/cuda/cuda_executor.cc:901] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero. See more at https://github.com/torvalds/linux/blob/v6.0/Documentation/ABI/testing/sysfs-bus-pci#L344-L355
2024-06-22 21:29:48.573876: I external/local_xla/xla/stream_executor/cuda/cuda_executor.cc:901] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero. See more at https://github.com/torvalds/linux/blob/v6.0/Documentation/ABI/testing/sysfs-bus-pci#L344-L355
2024-06-22 21:29:48.574048: I external/local_xla/xla/stream_executor/cuda/cuda_executor.cc:901] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero. See more at https://github.com/torvalds/linux/blob/v6.0/Documentation/ABI/testing/sysfs-bus-pci#L344-L355
2024-06-22 21:29:48.574145: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1929] Created device /job:localhost/replica:0/task:0/device:GPU:0 with 5070 MB memory: -> device: 0, name: NVIDIA GeForce GTX 1060 6GB, pci bus id: 0000:01:00.0, compute capability: 6.1
Found 10000 files belonging to 10 classes.
Found 10000 files belonging to 10 classes.
Found 10000 files belonging to 10 classes.
Found 10000 files belonging to 10 classes.
Found 10000 files belonging to 10 classes.
['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
```

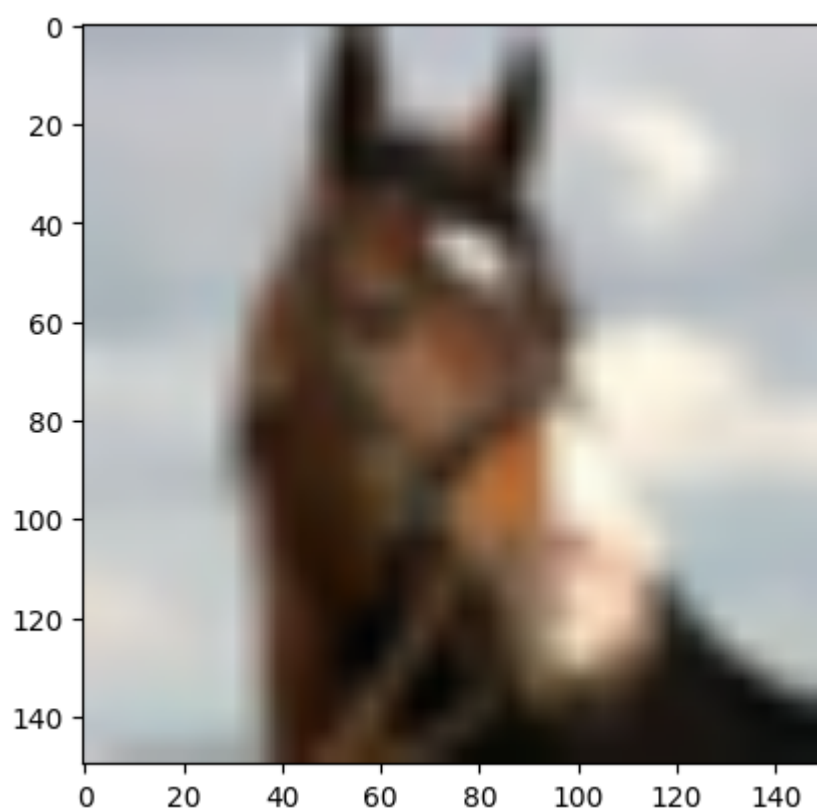
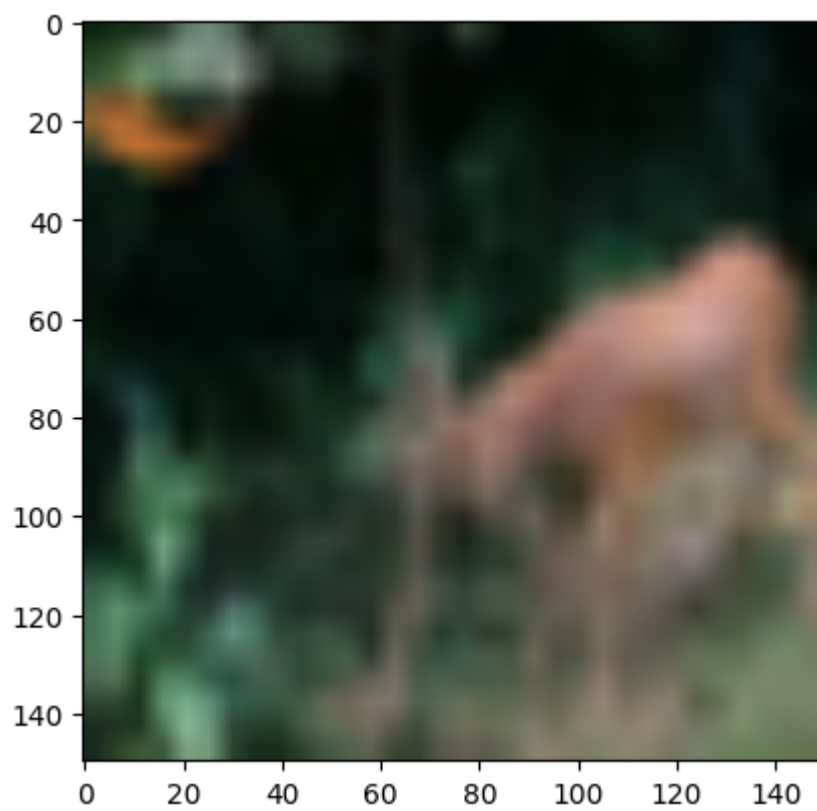
```
In [ ]: import matplotlib.pyplot as plt

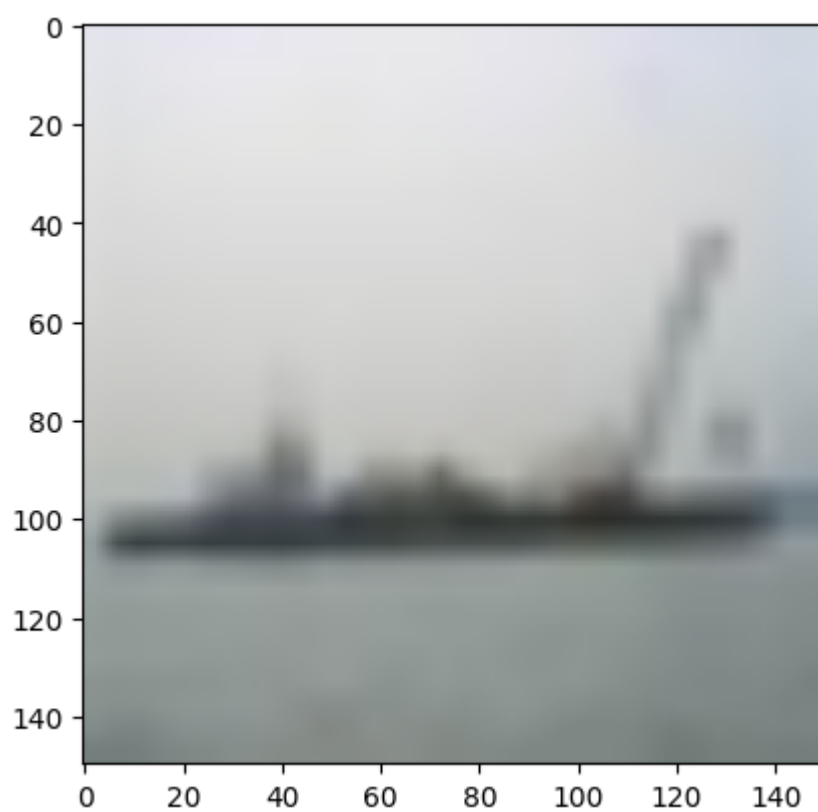
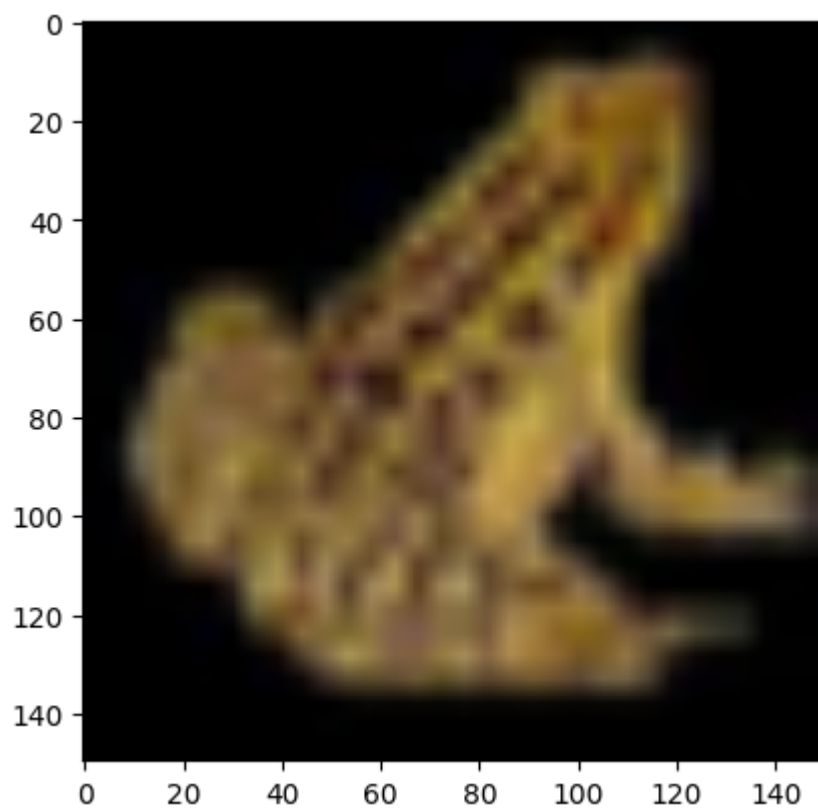
        for data, _ in train_dataset.take(1):
```

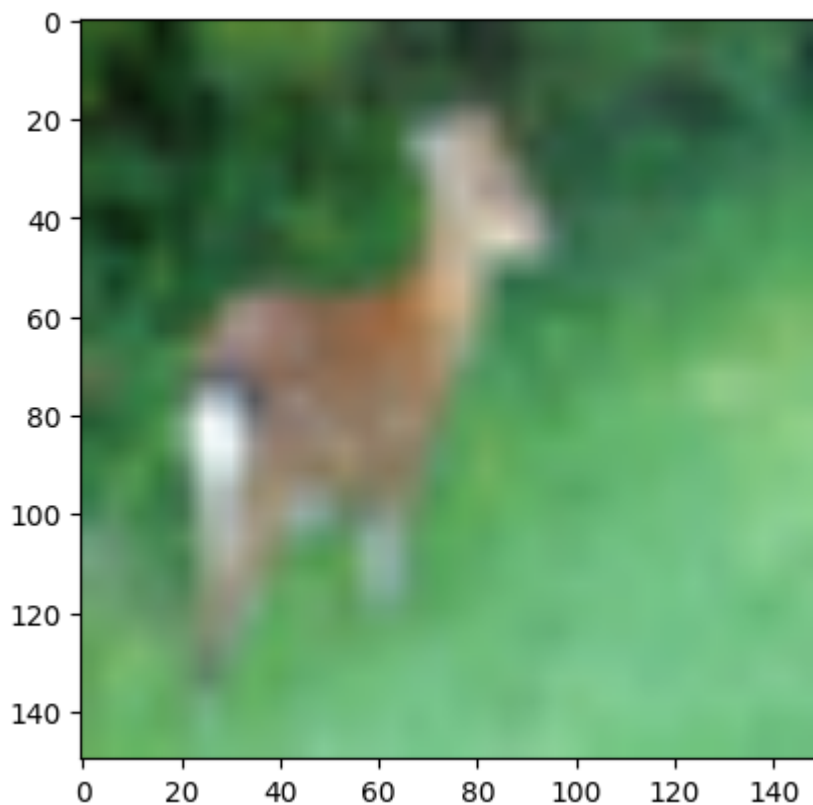
```
for i in range(9):  
    plt.imshow(data[i].numpy().astype('uint8'))  
    plt.show()  
break
```











## Feature extraction

Use the VGG16 pre-trained model to extract features from the images.

```
In [ ]: from tensorflow.keras.applications.vgg16 import VGG16

base_model = VGG16(include_top=False, weights='imagenet', input_shape=(IM
base_model.trainable = False # Freeze the base model
```

```
In [ ]: # from tensorflow import keras
# import numpy as np

# def get_features_and_labels(dataset):
#     all_features = []
#     all_labels = []
#     for images, labels in dataset:
#         preprocessed_images = keras.applications.vgg16.preprocess_input
#         features = base_model.predict(preprocessed_images)
#         all_features.append(features)
#         all_labels.append(labels)
#     return np.concatenate(all_features), np.concatenate(all_labels)

# validation_features, validation_labels = get_features_and_labels(valida
# np.save('validation_features.npy', validation_features)
# np.save('validation_labels.npy', validation_labels)
# validation_features = None
# validation_labels = None
# print("Validation features and labels saved")

# test_features, test_labels = get_features_and_labels(test_dataset)
# np.save('test_features.npy', test_features)
```



```
# np.save('test_labels.npy', test_labels)
# test_features = None
# test_labels = None
# print("Test features and labels saved")

# train_features, train_labels = get_features_and_labels(train_dataset)
# np.save('train_features.npy', train_features)
# np.save('train_labels.npy', train_labels)
# train_features = None
# train_labels = None
# print("Train features and labels saved")
```

## Load the extracted features

```
In [ ]: from numpy import load

train_features = load('features_T/train_features.npy')
train_labels = load('features_T/train_labels.npy')

validation_features = load('features_T/validation_features.npy')
validation_labels = load('features_T/validation_labels.npy')

test_features = load('features_T/test_features.npy')
test_labels = load('features_T/test_labels.npy')
```

# MODEL ARCHITECTURE

## Transfer Learning Model

Use the extracted features to train a model for the classification task.

Use two dense layers with 256 neurons and a in-between dropout layer with a rate of 0.5.

```
In [ ]: from tensorflow import keras
from keras.applications import VGG16
from keras import layers

inputs = keras.Input(shape=(4, 4, 512))

x = layers.Flatten()(inputs)
x = layers.Dense(256, activation="relu")(x)
x = layers.Dropout(0.5)(x)
x = layers.Dense(256, activation="relu")(x)
x = layers.Dropout(0.5)(x)

outputs = layers.Dense(10, activation="softmax")(x) # Softmax for multi-

model = keras.Model(inputs=inputs, outputs=outputs)
model.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
input_2 (InputLayer)	[(None, 4, 4, 512)]	0
flatten (Flatten)	(None, 8192)	0
dense (Dense)	(None, 256)	2097408
dropout (Dropout)	(None, 256)	0
dense_1 (Dense)	(None, 256)	65792
dropout_1 (Dropout)	(None, 256)	0
dense_2 (Dense)	(None, 10)	2570

---

=====  
Total params: 2165770 (8.26 MB)  
Trainable params: 2165770 (8.26 MB)  
Non-trainable params: 0 (0.00 Byte)  
=====

---

## Compile Model

### Loss function:

We use the *Categorical Crossentropy* loss function because it is a multi-class classification problem.

### Optimizer: Adam

We use the *Adam* optimizer because it is one of the best and most popular optimizers.

```
In [ ]: model.compile(  
        loss='categorical_crossentropy',  
        optimizer='adam',  
        metrics=['acc'])
```

## Train Model

Train the model with Early stopping, Model checkpoint, and Learning rate reduction callbacks.

```
In [ ]: from keras.callbacks import EarlyStopping, ModelCheckpoint, ReduceLRonPla  
  
learning_rate_reduction = ReduceLRonPlateau(  
    monitor='val_acc',  
    patience=3,  
    verbose=1,  
    factor=0.5,  
    min_lr=1e-6)  
  
early_stop = EarlyStopping(monitor='val_acc',  
                           patience=5,  
                           restore_best_weights=True)
```

```

model_checkpoint = ModelCheckpoint('models/T01/checkpoints/T01-cp.h5', sa

history = model.fit(
    train_features, train_labels,
    epochs=50,
    validation_data=(validation_features, validation_labels),
    callbacks=[early_stop, model_checkpoint, learning_rate_reduction])

```

2024-06-22 21:29:59.800833: W external/local\_tsl/tsl/framework/cpu\_allocator\_impl.cc:83] Allocation of 1310720000 exceeds 10% of free system memory.  
 2024-06-22 21:30:00.629393: W external/local\_tsl/tsl/framework/cpu\_allocator\_impl.cc:83] Allocation of 1310720000 exceeds 10% of free system memory.  
 Epoch 1/50

2024-06-22 21:30:02.762521: I external/local\_xla/xla/service/service.cc:168] XLA service 0x7db791188ae0 initialized for platform CUDA (this does not guarantee that XLA will be used). Devices:

2024-06-22 21:30:02.762543: I external/local\_xla/xla/service/service.cc:176] StreamExecutor device (0): NVIDIA GeForce GTX 1060 6GB, Compute Capability 6.1

2024-06-22 21:30:02.793368: I tensorflow/compiler/mlir/tensorflow/utils/dump\_mlir\_util.cc:269] disabling MLIR crash reproducer, set env var `MLIR\_CRASH\_REPRODUCER\_DIRECTORY` to enable.

2024-06-22 21:30:02.858548: I external/local\_xla/xla/stream\_executor/cuda/cuda\_dnn.cc:454] Loaded cuDNN version 8904

WARNING: All log messages before absl::InitializeLog() is called are written to STDERR

I0000 00:00:17.19088202.944203 4013 device\_compiler.h:186] Compiled cluster using XLA! This line is logged at most once for the lifetime of the process.

1250/1250 [=====] - 7s 4ms/step - loss: 1.6091 - acc: 0.5875 - val\_loss: 0.8074 - val\_acc: 0.7531 - lr: 0.0010

Epoch 2/50

20/1250 [.....] - ETA: 3s - loss: 0.9607 - acc: 0.6906

/home/pws/miniconda3/envs/tensorflow/lib/python3.11/site-packages/keras/src/engine/training.py:3103: UserWarning: You are saving your model as an HDF5 file via `model.save()`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my\_model.keras')`.

saving\_api.save\_model(

```
1250/1250 [=====] - 4s 4ms/step - loss: 0.9766 -  
acc: 0.7104 - val_loss: 0.5722 - val_acc: 0.8305 - lr: 0.0010  
Epoch 3/50  
1250/1250 [=====] - 5s 4ms/step - loss: 0.7366 -  
acc: 0.7757 - val_loss: 0.5336 - val_acc: 0.8462 - lr: 0.0010  
Epoch 4/50  
1250/1250 [=====] - 4s 3ms/step - loss: 0.6321 -  
acc: 0.8100 - val_loss: 0.4832 - val_acc: 0.8607 - lr: 0.0010  
Epoch 5/50  
1250/1250 [=====] - 5s 4ms/step - loss: 0.5985 -  
acc: 0.8208 - val_loss: 0.4522 - val_acc: 0.8618 - lr: 0.0010  
Epoch 6/50  
1250/1250 [=====] - 5s 4ms/step - loss: 0.5567 -  
acc: 0.8350 - val_loss: 0.4580 - val_acc: 0.8680 - lr: 0.0010  
Epoch 7/50  
1250/1250 [=====] - 4s 3ms/step - loss: 0.5391 -  
acc: 0.8418 - val_loss: 0.4585 - val_acc: 0.8613 - lr: 0.0010  
Epoch 8/50  
1250/1250 [=====] - 4s 3ms/step - loss: 0.4967 -  
acc: 0.8530 - val_loss: 0.4723 - val_acc: 0.8649 - lr: 0.0010  
Epoch 9/50  
1240/1250 [=====>.] - ETA: 0s - loss: 0.5026 - acc:  
0.8539  
Epoch 9: ReduceLROnPlateau reducing learning rate to 0.000500000023748725  
7.  
1250/1250 [=====] - 4s 3ms/step - loss: 0.5028 -  
acc: 0.8538 - val_loss: 0.4795 - val_acc: 0.8632 - lr: 0.0010  
Epoch 10/50  
1250/1250 [=====] - 4s 3ms/step - loss: 0.4069 -  
acc: 0.8728 - val_loss: 0.4460 - val_acc: 0.8729 - lr: 5.0000e-04  
Epoch 11/50  
1250/1250 [=====] - 4s 3ms/step - loss: 0.3741 -  
acc: 0.8830 - val_loss: 0.4377 - val_acc: 0.8785 - lr: 5.0000e-04  
Epoch 12/50  
1250/1250 [=====] - 4s 3ms/step - loss: 0.3625 -  
acc: 0.8867 - val_loss: 0.4363 - val_acc: 0.8715 - lr: 5.0000e-04  
Epoch 13/50  
1250/1250 [=====] - 4s 3ms/step - loss: 0.3295 -  
acc: 0.8942 - val_loss: 0.4495 - val_acc: 0.8750 - lr: 5.0000e-04  
Epoch 14/50  
1237/1250 [=====>.] - ETA: 0s - loss: 0.3203 - acc:  
0.8985  
Epoch 14: ReduceLROnPlateau reducing learning rate to 0.000250000011874362  
8.  
1250/1250 [=====] - 4s 3ms/step - loss: 0.3201 -  
acc: 0.8985 - val_loss: 0.4594 - val_acc: 0.8736 - lr: 5.0000e-04  
Epoch 15/50  
1250/1250 [=====] - 4s 3ms/step - loss: 0.2963 -  
acc: 0.9048 - val_loss: 0.4596 - val_acc: 0.8749 - lr: 2.5000e-04  
Epoch 16/50  
1250/1250 [=====] - 4s 3ms/step - loss: 0.2694 -  
acc: 0.9103 - val_loss: 0.4673 - val_acc: 0.8781 - lr: 2.5000e-04
```

## Save Model

```
In [ ]: keras.models.save_model(model, 'models/T01/T01-model.h5')
```

```
/tmp/ipykernel_3734/1570624933.py:1: UserWarning: You are saving your model as an HDF5 file via `model.save()`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')`.
keras.models.save_model(model, 'models/T01/T01-model.h5')
```

## Load Model

```
In [ ]: keras.models.load_model('models/T01/T01-model.h5')
```

```
Out[ ]: <keras.src.engine.functional.Functional at 0x7db834684110>
```

## EVALUATION

Evaluate the model on the validation dataset.

```
In [ ]: val_loss, val_acc = model.evaluate(validation_features, validation_labels)
print('val_acc:', val_acc)
```

```
313/313 [=====] - 1s 2ms/step - loss: 0.4377 - acc: 0.8785
val_acc: 0.8784999847412109
```

## Training and Validation Curves



Plot the training and validation accuracy and loss curves.

```
In [ ]: import matplotlib.pyplot as plt

# Extract the history from the training process
acc = history.history['acc']
val_acc = history.history['val_acc']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(acc) + 1)

# Plot the training and validation accuracy
plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()

# Plot the training and validation loss
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```



## Confusion Matrix

```
In [ ]: from sklearn.metrics import confusion_matrix
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```


y_pred = np.argmax(model.predict(validation_features), axis=1)
y_true = np.argmax(validation_labels, axis=1)

cm = confusion_matrix(y_true, y_pred)

plt.figure(figsize=(10, 8))
sns.heatmap(cm, annot=True, cmap='Blues', fmt='g')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()

```

313/313 [=====] - 0s 1ms/step



```

In [ ]: from sklearn.metrics import classification_report

report = classification_report(y_true, y_pred, target_names=class_names)
print(report)

```

	precision	recall	f1-score	support
airplane	0.91	0.90	0.90	1014
automobile	0.95	0.95	0.95	1014
bird	0.86	0.84	0.85	952
cat	0.68	0.81	0.74	1016
deer	0.89	0.83	0.86	997
dog	0.84	0.79	0.81	1025
frog	0.94	0.88	0.91	980
horse	0.88	0.90	0.89	977
ship	0.94	0.94	0.94	1003
truck	0.95	0.94	0.94	1022
accuracy			0.88	10000
macro avg	0.88	0.88	0.88	10000
weighted avg	0.88	0.88	0.88	10000

## Predictions

Predict and visualize the results for a sample image.

```
In [ ]: import tensorflow as tf
import matplotlib.pyplot as plt
from keras.preprocessing import image
from keras.applications.vgg16 import VGG16, preprocess_input
import numpy as np

# Load an image
img_path = train_dirs[0] + '/006_frog/alytes_obstetricans_s_000179.png'
img = tf.keras.preprocessing.image.load_img(img_path, target_size=(150, 150))

# Preprocess the image for VGG16
img_array = image.img_to_array(img)
img_array = np.expand_dims(img_array, axis=0)
img_array = preprocess_input(img_array)

plt.imshow(img)
plt.show()

print(img_array.shape)

base_model = VGG16(weights='imagenet', include_top=False, input_shape=(150, 150, 3))


# Extract features using VGG16
features = base_model.predict(img_array)

flattened_features = features.reshape((features.shape[0], -1))

# Predict using your custom model
result = model.predict(flattened_features)

print("Result: ", result.round())
print("Predicted class: ", class_names[np.argmax(result)])
print("True class: ", img_path.split('/')[-2].split('_')[-1])
```





```
(1, 150, 150, 3)
```

```
1/1 [=====] - 1s 1s/step
```

```
1/1 [=====] - 0s 67ms/step
```

```
Result: [[0. 0. 0. 0. 0. 0. 0. 1. 0. 0. 0.]]
```

```
Predicted class: frog
```

```
True class: frog
```