

Model_Transfer Learning 01 - With data augmentation - Feature extraction

Name: Alberto Pingo

Email: 2202145 @my.ipleiria.pt

Validation dataset: train5

Directories

This section sets up the directory paths used for training, validation, and test datasets based on the repository structure.

```
In [ ]: import os

current_dir = os.getcwd()

# TWO FOLDERS UP
data_dir = os.path.abspath(os.path.join(current_dir, os.pardir, os.pardir))
test_dir = os.path.join(data_dir, 'test')
train_dir = os.path.join(data_dir, 'train')

train_dirs = []
for i in range(1, 5):
    train_dirs.append(os.path.join(train_dir, 'train' + str(i)))

validation_dir = os.path.join(data_dir, 'train', 'train5')

print(current_dir)
print(data_dir)
print(test_dir)
print(train_dir)
print(validation_dir)
```

```
/home/pws/code/IA-image-classification/notebooks/models-T
/home/pws/code/IA-image-classification/data
/home/pws/code/IA-image-classification/data/test
/home/pws/code/IA-image-classification/data/train
/home/pws/code/IA-image-classification/data/train/train5
```

Preprocessing

Load the datasets and perform initial preprocessing. Images are resized to 32x32 pixels and batched.

```
In [ ]: from keras.utils import image_dataset_from_directory
import tensorflow as tf

# Load training datasets from train1 to train4
train_datasets = []
```

```

IMG_SIZE = 150
BATCH_SIZE = 32
train_dataset = image_dataset_from_directory(train_dirs[0], image_size=(I
# for i in range(1, 5):
#     dataset = image_dataset_from_directory(train_dirs[i-1], image_size=
#     train_datasets.append(dataset)

# train_dataset = train_datasets[0]
# for dataset in train_datasets[1:]:
#     train_dataset = train_dataset.concatenate(dataset)

# Load validation dataset
validation_dataset = image_dataset_from_directory(validation_dir, image_s

# Load test dataset
test_dataset = image_dataset_from_directory(test_dir, image_size=(IMG_SIZ

class_names = validation_dataset.class_names
class_names = [class_name.split('_')[-1] for class_name in class_names]

print(class_names)

```

Found 10000 files belonging to 10 classes.

Found 10000 files belonging to 10 classes.

Found 10000 files belonging to 10 classes.

['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse',
'ship', 'truck']

Configure the dataset for performance

```

In [ ]: AUTOTUNE = tf.data.AUTOTUNE

train_dataset = train_dataset.cache().shuffle(1000).prefetch(buffer_size=
validation_dataset = validation_dataset.cache().prefetch(buffer_size=AUTO
test_dataset = test_dataset.cache().prefetch(buffer_size=AUTOTUNE)

```

Data Augmentation

Random change of flipping the image horizontally.

Random chance of moving the image horizontally and vertically [-10%, 10%].

Tried with a more complex approach to data augmentation, but the results were worse because of the small size of the images.

```

In [ ]: from keras import layers

data_augmentation = tf.keras.Sequential([
    layers.RandomFlip("horizontal"),
    layers.RandomTranslation(0.1, 0.1, fill_mode='nearest'),
    layers.RandomRotation(0.03),
    layers.RandomZoom(0.03),
])

```

```

In [ ]: import matplotlib.pyplot as plt
import numpy as np

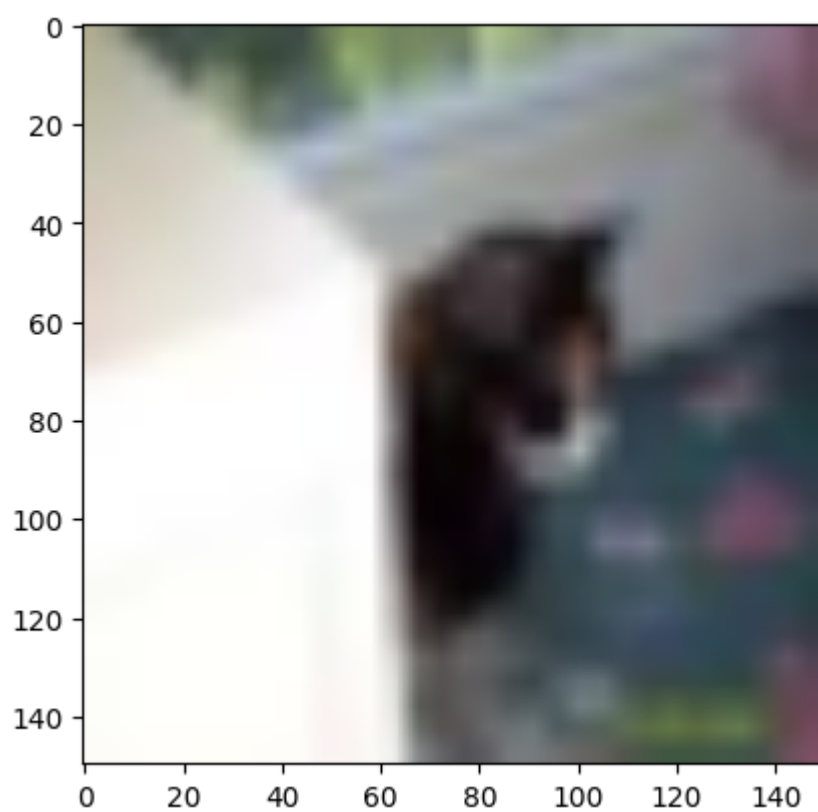
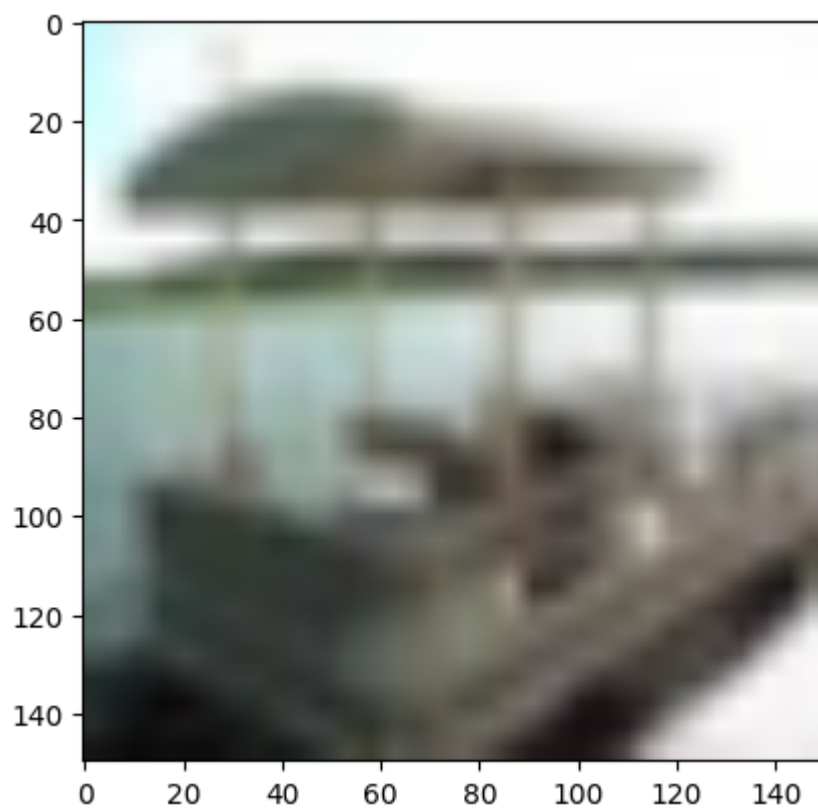
```

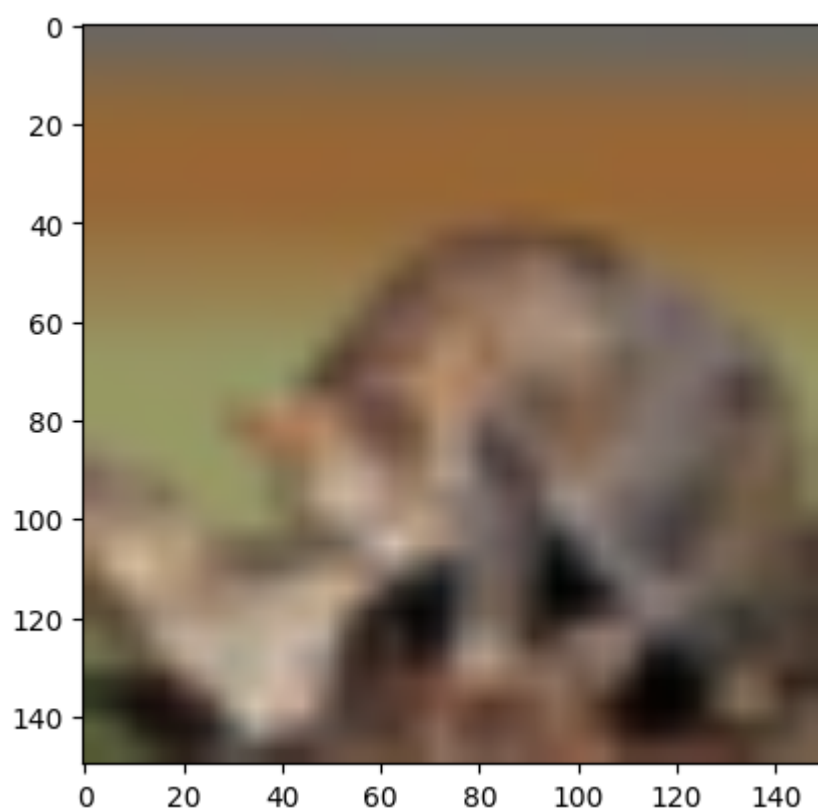
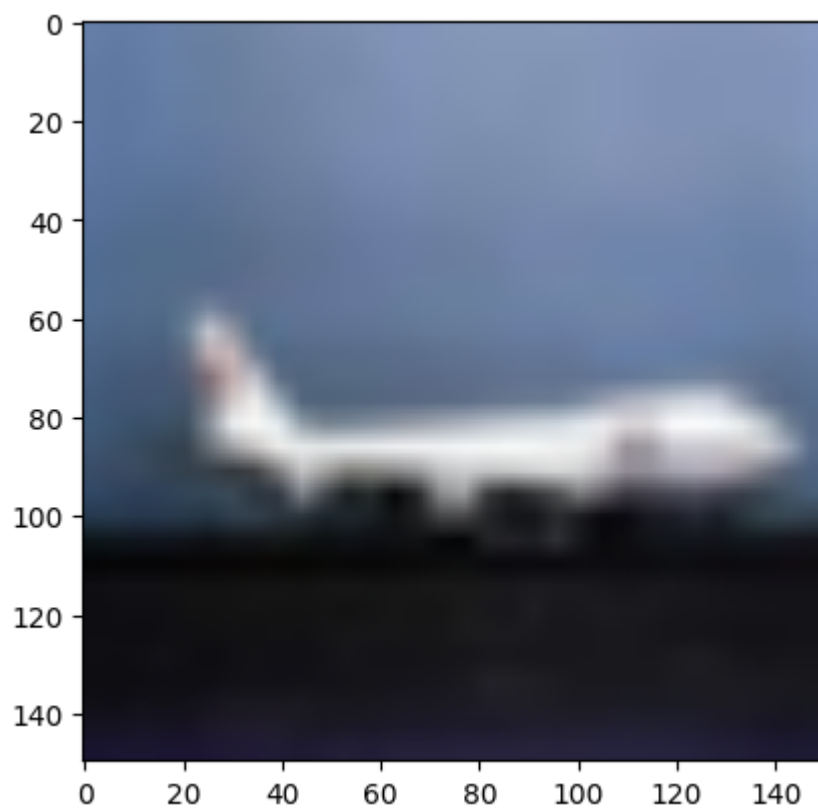
```
#Plot some Augmented images
for images, labels in train_dataset.take(1):
    plt.figure(figsize=(10, 10))
    first_image = images[0]
    for i in range(4):
        ax = plt.subplot(2, 2, i + 1)
        augmented_image = data_augmentation(tf.expand_dims(first_image, 0)
        plt.imshow(augmented_image[0] / 255)
        plt.axis('off')
```

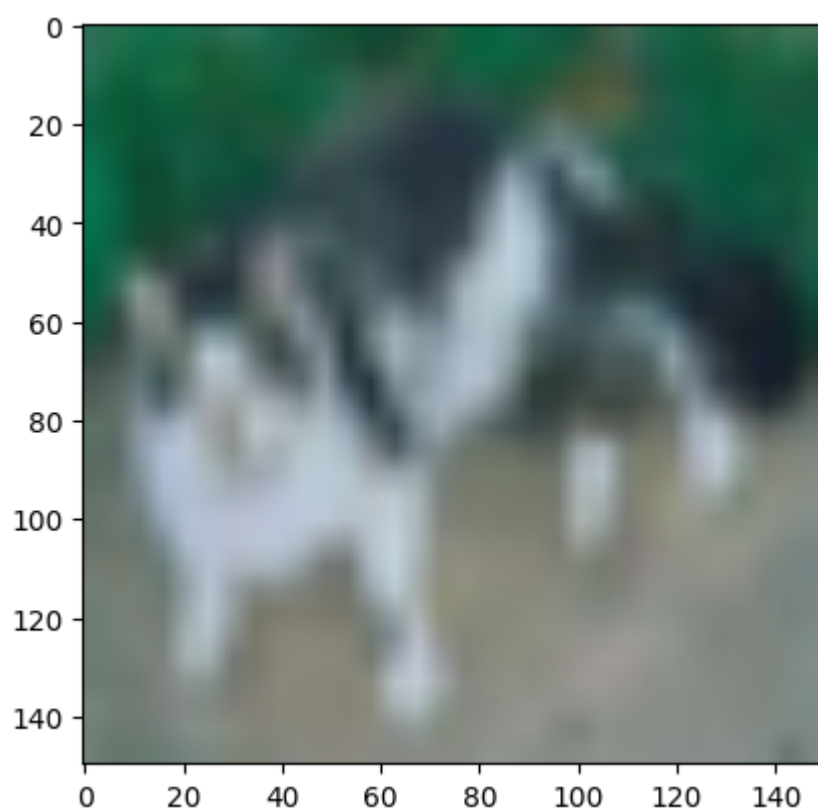
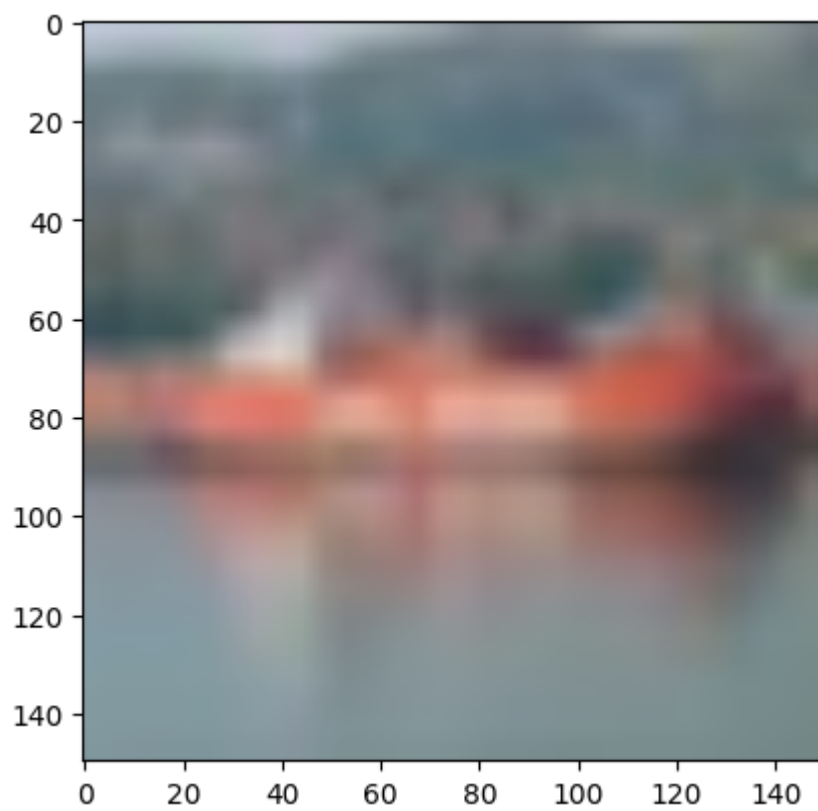


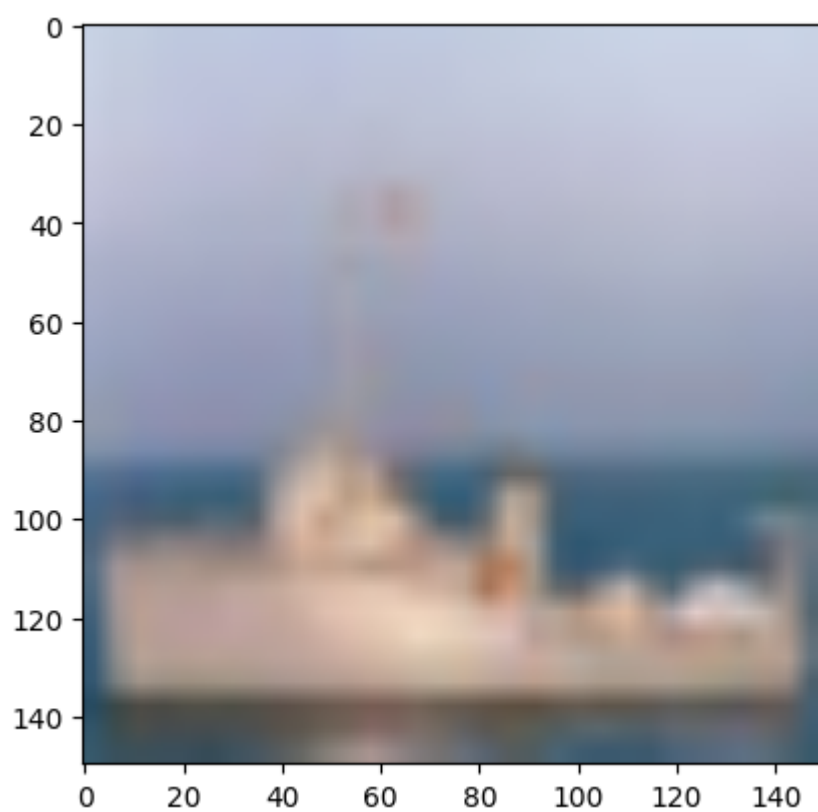
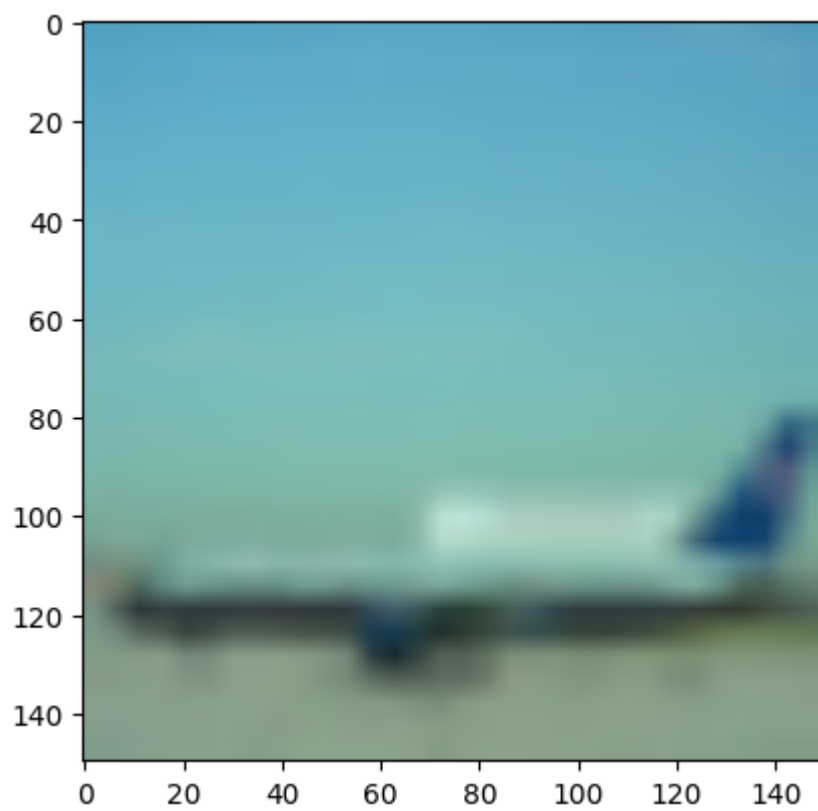
```
In [ ]: import matplotlib.pyplot as plt

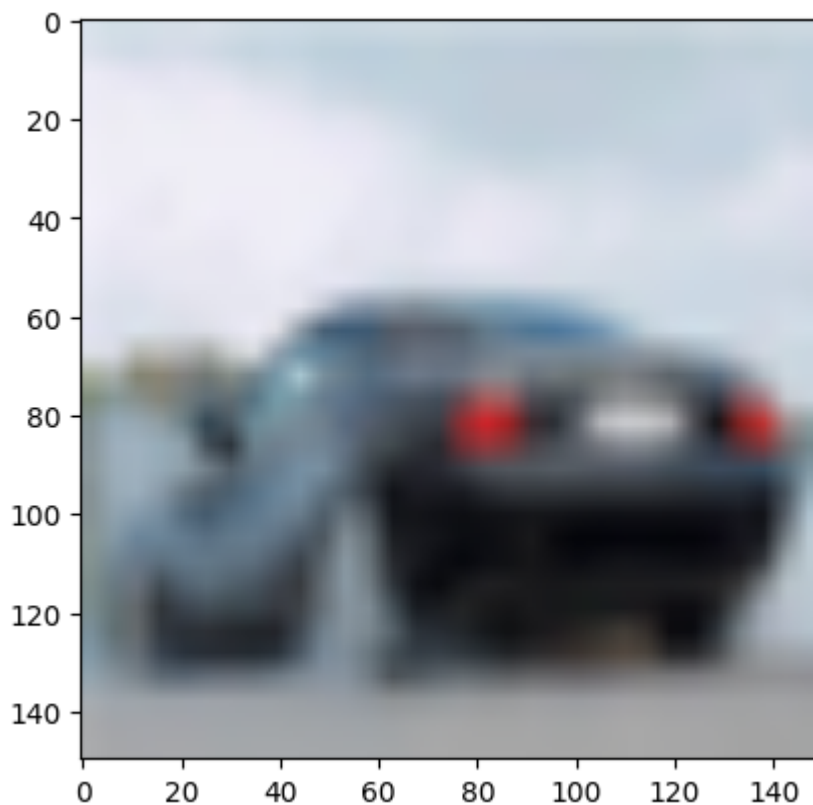
for data, _ in train_dataset.take(1):
    for i in range(9):
        plt.imshow(data[i].numpy().astype('uint8'))
        plt.show()
    break
```











MODEL ARCHITECTURE

Build a Convolutional Neural Network (CNN) model.

Transfer Learning Model

Use the VGG16 model as a base model.

Freeze the convolutional base and train the densely connected classifier.

Data augmentation before processing the images to the model.

Use two dense layers with 256 neurons and a in-between dropout layer with a rate of 0.5.

```
In [ ]: from tensorflow import keras
        from keras.applications.vgg16 import VGG16
        from keras import layers

        base_model = VGG16(include_top=False, weights='imagenet')
        base_model.trainable = False # Freeze the base model

        inputs = keras.Input(shape=(IMG_SIZE, IMG_SIZE, 3))
        x = data_augmentation(inputs)
        x = keras.applications.vgg16.preprocess_input(x)
        x = base_model(x)
```



```
x = layers.Flatten()(x)
x = layers.Dense(256, activation="relu")(x)
x = layers.Dropout(0.5)(x)
x = layers.Dense(256, activation="relu")(x)
x = layers.Dropout(0.5)(x)

outputs = layers.Dense(10, activation="softmax")(x) # Softmax for multi-

model = keras.Model(inputs=inputs, outputs=outputs)
model.summary()
```

Model: "model_1"

Layer (type)	Output Shape	Param #
input_5 (InputLayer)	[(None, 150, 150, 3)]	0
sequential_1 (Sequential)	(None, 150, 150, 3)	0
tf.__operators__.getitem_1 (SlicingOpLambda)	(None, 150, 150, 3)	0
tf.nn.bias_add_1 (TFOpLambda)	(None, 150, 150, 3)	0
vgg16 (Functional)	(None, None, None, 512)	14714688
flatten_1 (Flatten)	(None, 8192)	0
dense_3 (Dense)	(None, 256)	2097408
dropout_2 (Dropout)	(None, 256)	0
dense_4 (Dense)	(None, 256)	65792
dropout_3 (Dropout)	(None, 256)	0
dense_5 (Dense)	(None, 10)	2570

Layer (type)	Output Shape	Param #
input_5 (InputLayer)	[(None, 150, 150, 3)]	0
sequential_1 (Sequential)	(None, 150, 150, 3)	0
tf.__operators__.getitem_1 (SlicingOpLambda)	(None, 150, 150, 3)	0
tf.nn.bias_add_1 (TFOpLambda)	(None, 150, 150, 3)	0
vgg16 (Functional)	(None, None, None, 512)	14714688
flatten_1 (Flatten)	(None, 8192)	0
dense_3 (Dense)	(None, 256)	2097408
dropout_2 (Dropout)	(None, 256)	0
dense_4 (Dense)	(None, 256)	65792
dropout_3 (Dropout)	(None, 256)	0
dense_5 (Dense)	(None, 10)	2570

Total params: 16880458 (64.39 MB)
Trainable params: 2165770 (8.26 MB)
Non-trainable params: 14714688 (56.13 MB)

Compile Model

Loss function:

We use the *Categorical Crossentropy* loss function because it is a multi-class classification problem.

Optimizer: Adam

We use the *Adam* optimizer because it is one of the best and most popular optimizers.

```
In [ ]: model.compile(  
        loss='categorical_crossentropy',  
        optimizer='adam',  
        metrics=['acc'])
```

Train Model

Train the model with Early stopping, Model checkpoint, and Learning rate reduction callbacks.

```
In [ ]: from keras.callbacks import EarlyStopping, ModelCheckpoint, ReduceLROnPla  
  
learning_rate_reduction = ReduceLROnPlateau(  
    monitor='val_acc',  
    patience=3,  
    verbose=1,  
    factor=0.4,  
    min_lr=1e-6)  
  
early_stop = EarlyStopping(monitor='val_acc',  
                           patience=8,  
                           restore_best_weights=True)  
model_checkpoint = ModelCheckpoint('models/T01/checkpoints/T01-DA-cp.h5',  
  
history = model.fit(  
    train_dataset,  
    epochs=50,  
    validation_data=validation_dataset,  
    callbacks=[early_stop, model_checkpoint, learning_rate_reduction])
```

Epoch 1/50

313/313 [=====] - ETA: 0s - loss: 2.6981 - acc: 0.4071

/home/pws/miniconda3/envs/tensorflow/lib/python3.11/site-packages/keras/src/engine/training.py:3103: UserWarning: You are saving your model as an HDF5 file via `model.save()`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')`.

saving_api.save_model(
 saving_api.save_model(
 model, 'models/T01/checkpoints/T01-DA-cp.h5',
 overwrite=True)

```
313/313 [=====] - 57s 177ms/step - loss: 2.6981 -  
acc: 0.4071 - val_loss: 1.0745 - val_acc: 0.6791 - lr: 0.0010  
Epoch 2/50  
Epoch 2/50  
313/313 [=====] - 55s 177ms/step - loss: 1.5151 -  
acc: 0.5776 - val_loss: 0.8892 - val_acc: 0.7617 - lr: 0.0010  
Epoch 3/50  
313/313 [=====] - 53s 169ms/step - loss: 1.3404 -  
acc: 0.6067 - val_loss: 0.8123 - val_acc: 0.7691 - lr: 0.0010  
Epoch 4/50  
313/313 [=====] - 53s 169ms/step - loss: 1.2438 -  
acc: 0.6428 - val_loss: 0.7669 - val_acc: 0.7866 - lr: 0.0010  
Epoch 5/50  
313/313 [=====] - 53s 170ms/step - loss: 1.1292 -  
acc: 0.6721 - val_loss: 0.7313 - val_acc: 0.7917 - lr: 0.0010  
Epoch 6/50  
313/313 [=====] - 52s 167ms/step - loss: 1.0273 -  
acc: 0.6853 - val_loss: 0.6603 - val_acc: 0.8050 - lr: 0.0010  
Epoch 7/50  
313/313 [=====] - 52s 167ms/step - loss: 0.9732 -  
acc: 0.7094 - val_loss: 0.6268 - val_acc: 0.8115 - lr: 0.0010  
Epoch 8/50  
313/313 [=====] - 52s 167ms/step - loss: 0.8859 -  
acc: 0.7272 - val_loss: 0.6241 - val_acc: 0.8113 - lr: 0.0010  
Epoch 9/50  
313/313 [=====] - 52s 167ms/step - loss: 0.8311 -  
acc: 0.7493 - val_loss: 0.5691 - val_acc: 0.8303 - lr: 0.0010  
Epoch 10/50  
313/313 [=====] - 52s 166ms/step - loss: 0.7787 -  
acc: 0.7570 - val_loss: 0.5724 - val_acc: 0.8255 - lr: 0.0010  
Epoch 11/50  
313/313 [=====] - 53s 169ms/step - loss: 0.7659 -  
acc: 0.7566 - val_loss: 0.5308 - val_acc: 0.8362 - lr: 0.0010  
Epoch 12/50  
313/313 [=====] - 54s 171ms/step - loss: 0.7219 -  
acc: 0.7757 - val_loss: 0.5357 - val_acc: 0.8373 - lr: 0.0010  
Epoch 13/50  
313/313 [=====] - 55s 175ms/step - loss: 0.7197 -  
acc: 0.7723 - val_loss: 0.5222 - val_acc: 0.8393 - lr: 0.0010  
Epoch 14/50  
313/313 [=====] - 55s 175ms/step - loss: 0.6872 -  
acc: 0.7885 - val_loss: 0.5169 - val_acc: 0.8463 - lr: 0.0010  
Epoch 15/50  
313/313 [=====] - 54s 174ms/step - loss: 0.6878 -  
acc: 0.7906 - val_loss: 0.5759 - val_acc: 0.8322 - lr: 0.0010  
Epoch 16/50  
313/313 [=====] - 55s 175ms/step - loss: 0.6585 -  
acc: 0.7965 - val_loss: 0.5346 - val_acc: 0.8425 - lr: 0.0010  
Epoch 17/50  
313/313 [=====] - ETA: 0s - loss: 0.6559 - acc:  
0.8011  
Epoch 17: ReduceLROnPlateau reducing learning rate to 0.000400000018998980  
5.  
313/313 [=====] - 54s 174ms/step - loss: 0.6559 -  
acc: 0.8011 - val_loss: 0.5280 - val_acc: 0.8399 - lr: 0.0010  
Epoch 18/50  
313/313 [=====] - 55s 175ms/step - loss: 0.5827 -  
acc: 0.8172 - val_loss: 0.4923 - val_acc: 0.8531 - lr: 4.0000e-04  
Epoch 19/50  
313/313 [=====] - 54s 174ms/step - loss: 0.5292 -
```

```
acc: 0.8283 - val_loss: 0.4893 - val_acc: 0.8537 - lr: 4.0000e-04
Epoch 20/50
313/313 [=====] - 54s 173ms/step - loss: 0.5384 -
acc: 0.8275 - val_loss: 0.4746 - val_acc: 0.8543 - lr: 4.0000e-04
Epoch 21/50
313/313 [=====] - 54s 174ms/step - loss: 0.5044 -
acc: 0.8377 - val_loss: 0.4802 - val_acc: 0.8526 - lr: 4.0000e-04
Epoch 22/50
313/313 [=====] - 55s 176ms/step - loss: 0.4972 -
acc: 0.8363 - val_loss: 0.4868 - val_acc: 0.8523 - lr: 4.0000e-04
Epoch 23/50
313/313 [=====] - 55s 176ms/step - loss: 0.4787 -
acc: 0.8432 - val_loss: 0.4841 - val_acc: 0.8545 - lr: 4.0000e-04
Epoch 24/50
313/313 [=====] - 56s 178ms/step - loss: 0.4606 -
acc: 0.8497 - val_loss: 0.4708 - val_acc: 0.8567 - lr: 4.0000e-04
Epoch 25/50
313/313 [=====] - 56s 179ms/step - loss: 0.4757 -
acc: 0.8498 - val_loss: 0.4629 - val_acc: 0.8602 - lr: 4.0000e-04
Epoch 26/50
313/313 [=====] - 55s 176ms/step - loss: 0.4619 -
acc: 0.8522 - val_loss: 0.4729 - val_acc: 0.8572 - lr: 4.0000e-04
Epoch 27/50
313/313 [=====] - 55s 176ms/step - loss: 0.4767 -
acc: 0.8495 - val_loss: 0.4731 - val_acc: 0.8534 - lr: 4.0000e-04
Epoch 28/50
313/313 [=====] - ETA: 0s - loss: 0.4416 - acc:
0.8557
Epoch 28: ReduceLRonPlateau reducing learning rate to 0.000160000007599592
22.
313/313 [=====] - 54s 174ms/step - loss: 0.4416 -
acc: 0.8557 - val_loss: 0.4616 - val_acc: 0.8587 - lr: 4.0000e-04
Epoch 29/50
313/313 [=====] - 55s 175ms/step - loss: 0.4336 -
acc: 0.8566 - val_loss: 0.4572 - val_acc: 0.8590 - lr: 1.6000e-04
Epoch 30/50
313/313 [=====] - 55s 176ms/step - loss: 0.4275 -
acc: 0.8613 - val_loss: 0.4562 - val_acc: 0.8589 - lr: 1.6000e-04
Epoch 31/50
313/313 [=====] - ETA: 0s - loss: 0.4054 - acc:
0.8664
Epoch 31: ReduceLRonPlateau reducing learning rate to 6.40000042039901e-0
5.
313/313 [=====] - 56s 178ms/step - loss: 0.4054 -
acc: 0.8664 - val_loss: 0.4561 - val_acc: 0.8579 - lr: 1.6000e-04
Epoch 32/50
313/313 [=====] - 55s 176ms/step - loss: 0.3873 -
acc: 0.8725 - val_loss: 0.4518 - val_acc: 0.8593 - lr: 6.4000e-05
Epoch 33/50
313/313 [=====] - 54s 174ms/step - loss: 0.3928 -
acc: 0.8714 - val_loss: 0.4520 - val_acc: 0.8609 - lr: 6.4000e-05
Epoch 34/50
313/313 [=====] - 55s 175ms/step - loss: 0.3820 -
acc: 0.8736 - val_loss: 0.4516 - val_acc: 0.8598 - lr: 6.4000e-05
Epoch 35/50
313/313 [=====] - 54s 174ms/step - loss: 0.3774 -
acc: 0.8736 - val_loss: 0.4491 - val_acc: 0.8603 - lr: 6.4000e-05
Epoch 36/50
313/313 [=====] - 55s 175ms/step - loss: 0.3745 -
acc: 0.8753 - val_loss: 0.4477 - val_acc: 0.8612 - lr: 6.4000e-05
```

```
Epoch 37/50
113/113 [=====] - 54s 174ms/step - loss: 0.3749 -
acc: 0.8784 - val_loss: 0.4501 - val_acc: 0.8598 - lr: 6.4000e-05
Epoch 38/50
113/113 [=====] - 54s 174ms/step - loss: 0.3760 -
acc: 0.8762 - val_loss: 0.4480 - val_acc: 0.8603 - lr: 6.4000e-05
Epoch 39/50
113/113 [=====] - ETA: 0s - loss: 0.3626 - acc:
0.8803
Epoch 39: ReduceLR0nPlateau reducing learning rate to 2.560000284574926e-0
5.
113/113 [=====] - 55s 175ms/step - loss: 0.3626 -
acc: 0.8803 - val_loss: 0.4473 - val_acc: 0.8602 - lr: 6.4000e-05
Epoch 40/50
113/113 [=====] - 55s 175ms/step - loss: 0.3575 -
acc: 0.8765 - val_loss: 0.4459 - val_acc: 0.8617 - lr: 2.5600e-05
Epoch 41/50
113/113 [=====] - 54s 174ms/step - loss: 0.3672 -
acc: 0.8768 - val_loss: 0.4460 - val_acc: 0.8606 - lr: 2.5600e-05
Epoch 42/50
113/113 [=====] - 55s 175ms/step - loss: 0.3639 -
acc: 0.8814 - val_loss: 0.4457 - val_acc: 0.8610 - lr: 2.5600e-05
Epoch 43/50
113/113 [=====] - ETA: 0s - loss: 0.3672 - acc:
0.8825
Epoch 43: ReduceLR0nPlateau reducing learning rate to 1.0240000847261399
e-05.
113/113 [=====] - 55s 175ms/step - loss: 0.3672 -
acc: 0.8825 - val_loss: 0.4442 - val_acc: 0.8614 - lr: 2.5600e-05
Epoch 44/50
113/113 [=====] - 55s 175ms/step - loss: 0.3639 -
acc: 0.8827 - val_loss: 0.4437 - val_acc: 0.8618 - lr: 1.0240e-05
Epoch 45/50
113/113 [=====] - 54s 173ms/step - loss: 0.3661 -
acc: 0.8756 - val_loss: 0.4441 - val_acc: 0.8617 - lr: 1.0240e-05
Epoch 46/50
113/113 [=====] - 60s 191ms/step - loss: 0.3566 -
acc: 0.8847 - val_loss: 0.4446 - val_acc: 0.8615 - lr: 1.0240e-05
Epoch 47/50
113/113 [=====] - ETA: 0s - loss: 0.3644 - acc:
0.8796
Epoch 47: ReduceLR0nPlateau reducing learning rate to 4.09600033890456e-0
6.
113/113 [=====] - 59s 188ms/step - loss: 0.3644 -
acc: 0.8796 - val_loss: 0.4451 - val_acc: 0.8607 - lr: 1.0240e-05
Epoch 48/50
113/113 [=====] - 59s 190ms/step - loss: 0.3515 -
acc: 0.8854 - val_loss: 0.4449 - val_acc: 0.8605 - lr: 4.0960e-06
Epoch 49/50
113/113 [=====] - 59s 189ms/step - loss: 0.3617 -
acc: 0.8823 - val_loss: 0.4451 - val_acc: 0.8604 - lr: 4.0960e-06
Epoch 50/50
113/113 [=====] - ETA: 0s - loss: 0.3562 - acc:
0.8784
Epoch 50: ReduceLR0nPlateau reducing learning rate to 1.6384001355618238
e-06.
113/113 [=====] - 59s 189ms/step - loss: 0.3562 -
acc: 0.8784 - val_loss: 0.4452 - val_acc: 0.8609 - lr: 4.0960e-06
```

Save Model

```
In [ ]: keras.models.save_model(model, 'models/T01/T01-DA-model.h5')
```

```
/tmp/ipykernel_8035/1849881407.py:1: UserWarning: You are saving your model as an HDF5 file via `model.save()`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')`.
```

```
keras.models.save_model(model, 'models/T01/T01-DA-model.h5')
```

```
In [ ]: keras.models.load_model('models/T01/T01-DA-model.h5')
```

```
Out[ ]: <keras.src.engine.functional.Functional at 0x72a4b2d92ad0>
```

EVALUATION

Evaluate the model on the validation dataset.

```
In [ ]: val_loss, val_acc = model.evaluate(validation_dataset)
print('val_acc:', val_acc)
```

```
313/313 [=====] - 29s 93ms/step - loss: 0.4452 - acc: 0.8609
val_acc: 0.8608999848365784
```

Training and Validation Curves



Plot the training and validation accuracy and loss curves.

```
In [ ]: import matplotlib.pyplot as plt

# Extract the history from the training process
acc = history.history['acc']
val_acc = history.history['val_acc']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(acc) + 1)

# Plot the training and validation accuracy
plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()

# Plot the training and validation loss
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```



Confusion Matrix

```
In [ ]: from sklearn.metrics import confusion_matrix
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```



```
y_true = []
y_pred = []

for images, labels in validation_dataset:
    y_true.extend(np.argmax(labels, axis=1))
    y_pred.extend(np.argmax(model.predict(images), axis=1))

y_true = np.array(y_true)
y_pred = np.array(y_pred)

cm = confusion_matrix(y_true, y_pred)

plt.figure(figsize=(10, 8))
sns.heatmap(cm, annot=True, cmap='Blues', fmt='g')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()
```

```
1/1 [=====] - 0s 162ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 19ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 20ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 20ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 20ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 20ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 19ms/step
1/1 [=====] - 0s 20ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 20ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 19ms/step
```

1/1	[=====]	- 0s	16ms/step
1/1	[=====]	- 0s	20ms/step
1/1	[=====]	- 0s	17ms/step
1/1	[=====]	- 0s	38ms/step
1/1	[=====]	- 0s	16ms/step
1/1	[=====]	- 0s	16ms/step
1/1	[=====]	- 0s	18ms/step
1/1	[=====]	- 0s	16ms/step
1/1	[=====]	- 0s	16ms/step
1/1	[=====]	- 0s	19ms/step
1/1	[=====]	- 0s	18ms/step
1/1	[=====]	- 0s	16ms/step
1/1	[=====]	- 0s	16ms/step
1/1	[=====]	- 0s	17ms/step
1/1	[=====]	- 0s	21ms/step
1/1	[=====]	- 0s	19ms/step
1/1	[=====]	- 0s	17ms/step
1/1	[=====]	- 0s	16ms/step
1/1	[=====]	- 0s	17ms/step
1/1	[=====]	- 0s	18ms/step
1/1	[=====]	- 0s	16ms/step
1/1	[=====]	- 0s	22ms/step
1/1	[=====]	- 0s	17ms/step
1/1	[=====]	- 0s	16ms/step
1/1	[=====]	- 0s	16ms/step
1/1	[=====]	- 0s	16ms/step
1/1	[=====]	- 0s	24ms/step
1/1	[=====]	- 0s	18ms/step
1/1	[=====]	- 0s	16ms/step
1/1	[=====]	- 0s	17ms/step
1/1	[=====]	- 0s	16ms/step
1/1	[=====]	- 0s	16ms/step
1/1	[=====]	- 0s	16ms/step
1/1	[=====]	- 0s	17ms/step
1/1	[=====]	- 0s	17ms/step
1/1	[=====]	- 0s	18ms/step
1/1	[=====]	- 0s	20ms/step
1/1	[=====]	- 0s	17ms/step
1/1	[=====]	- 0s	17ms/step
1/1	[=====]	- 0s	17ms/step
1/1	[=====]	- 0s	16ms/step
1/1	[=====]	- 0s	16ms/step
1/1	[=====]	- 0s	17ms/step
1/1	[=====]	- 0s	16ms/step
1/1	[=====]	- 0s	17ms/step
1/1	[=====]	- 0s	18ms/step
1/1	[=====]	- 0s	18ms/step
1/1	[=====]	- 0s	16ms/step
1/1	[=====]	- 0s	17ms/step
1/1	[=====]	- 0s	17ms/step
1/1	[=====]	- 0s	17ms/step
1/1	[=====]	- 0s	17ms/step
1/1	[=====]	- 0s	18ms/step
1/1	[=====]	- 0s	16ms/step
1/1	[=====]	- 0s	17ms/step
1/1	[=====]	- 0s	17ms/step
1/1	[=====]	- 0s	18ms/step
1/1	[=====]	- 0s	16ms/step
1/1	[=====]	- 0s	19ms/step
1/1	[=====]	- 0s	17ms/step
1/1	[=====]	- 0s	17ms/step

```
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 19ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 19ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 19ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 19ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 43ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 18ms/step
```


```
1/1 [=====] - 0s 20ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 20ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 20ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
```

```
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 20ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 19ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 19ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
```

```

1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 15ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 16ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 132ms/step

```



```

In [ ]: print("Confusion Matrix:")
        print(cm)

```

```

Confusion Matrix:
[[876 11 17 10 18 2 4 9 47 20]
 [ 0 929 2 2 3 0 0 2 9 67]
 [ 24 0 771 35 75 12 23 8 2 2]
 [ 5 4 28 767 42 96 41 20 9 4]
 [ 18 0 30 19 846 14 30 35 3 2]
 [ 2 1 13 138 48 773 13 35 0 2]
 [ 3 0 25 26 22 7 891 1 3 2]
 [ 2 1 8 31 75 24 3 826 2 5]
 [ 15 8 5 3 2 0 3 2 953 12]
 [ 10 15 0 4 2 0 1 0 13 977]]

```

```
In [ ]: from sklearn.metrics import classification_report

report = classification_report(y_true, y_pred, target_names=class_names)
print(report)
```

	precision	recall	f1-score	support
airplane	0.92	0.86	0.89	1014
automobile	0.96	0.92	0.94	1014
bird	0.86	0.81	0.83	952
cat	0.74	0.75	0.75	1016
deer	0.75	0.85	0.79	997
dog	0.83	0.75	0.79	1025
frog	0.88	0.91	0.90	980
horse	0.88	0.85	0.86	977
ship	0.92	0.95	0.93	1003
truck	0.89	0.96	0.92	1022
accuracy			0.86	10000
macro avg	0.86	0.86	0.86	10000
weighted avg	0.86	0.86	0.86	10000

Predictions

Predict and visualize the results for a sample image.


```
In [ ]: import tensorflow as tf
import matplotlib.pyplot as plt
from keras.preprocessing import image

# Load an image
img_path = train_dirs[0] + '/006_frog/alytes_obstetricans_s_000179.png'
# img_path2 = train_dirs[0] + '/000_airplane/airbus_s_000012.png'
img = tf.keras.preprocessing.image.load_img(img_path, target_size=(150, 150))

# Preprocess the image
img_array = image.img_to_array(img)
img_array = tf.expand_dims(img_array, 0)

plt.imshow(img)
plt.show()

print(img_array.shape)
result = model.predict(img_array)
print("Result: ", result.round())
print("Predicted class: ", class_names[np.argmax(result)])
print("True class: ", img_path.split('/')[-2].split('_')[-1])
```

(1, 150, 150, 3)

1/1 [=====] - 0s 29ms/step

Result: [[0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]]

Predicted class: frog

True class: frog